



# KUBERNETES HANDBOOK

From Beginner to Expert

For Cloud & DevOps Engineers

By :

SANDIP DAS  
[LearnXOps.com](http://LearnXOps.com)

# Table of Contents

01

Introduction

02

Set-up Environment

03

Core Concepts in Action

04

Advanced Workloads and Patterns

05

Networking Deep Dive

06

Storage Solutions

07

Troubleshooting and Error Handling

08

Optimization & Cost Management

09

Kubernetes Gotchas and Anti-Patterns

10

Kubernetes Security Best Practices

11

Monitoring, Logging, and Observability

12

CI/CD Pipelines with Kubernetes

13

What Next You May Explore?

14

GET SET GO



# ABOUT ME



Hi! I'm Sandip Das, a Senior Cloud, DevOps & MLOps Engineer with 12+ years of experience, primarily working on production-level Kubernetes projects. Around 70-80% of my daily work involves Kubernetes, focusing on automation, scaling, and optimizing workloads for clients across the US, UK, EU, and UAE.

# Introduction

## What is Kubernetes ?

Kubernetes is an **open-source container orchestration platform** that automates deployment, scaling, and management of containerized applications. It enables efficient workload distribution, self-healing, and service discovery across a cluster of nodes.



## Why use Kubernetes?

- **Scalability:** Automatically scales applications horizontally and vertically based on demand, ensuring efficient handling of dynamic workloads.
- **High Availability:** Ensures applications stay up with self-healing capabilities, automatic rescheduling, and failover mechanisms.
- **Efficient Resource Utilization:** Optimizes container scheduling and resource allocation, ensuring cost-effectiveness and performance.
- **Automation:** Handles deployment, updates, and rollbacks seamlessly, reducing manual intervention and human errors.
- **Portability:** Runs across different cloud providers and on-premises environments, ensuring flexibility and avoiding vendor lock-in.
- **Service Discovery & Load Balancing:** Distributes traffic efficiently to services, ensuring optimal performance and reliability.
- **Security & Isolation:** Supports Role-Based Access Control (RBAC), namespaces, and network policies for secure, multi-tenant environments.
- **Extensibility:** Easily integrates with CI/CD pipelines, monitoring tools, and custom controllers, enabling advanced automation and observability.
- **Microservices Management:** Efficiently orchestrates decentralized services with automated deployment, service discovery, and load balancing.
- **Cloud-Native Ecosystem:** Seamlessly integrates with cloud services, DevOps tools, and modern application development workflows.



@Sandip Das

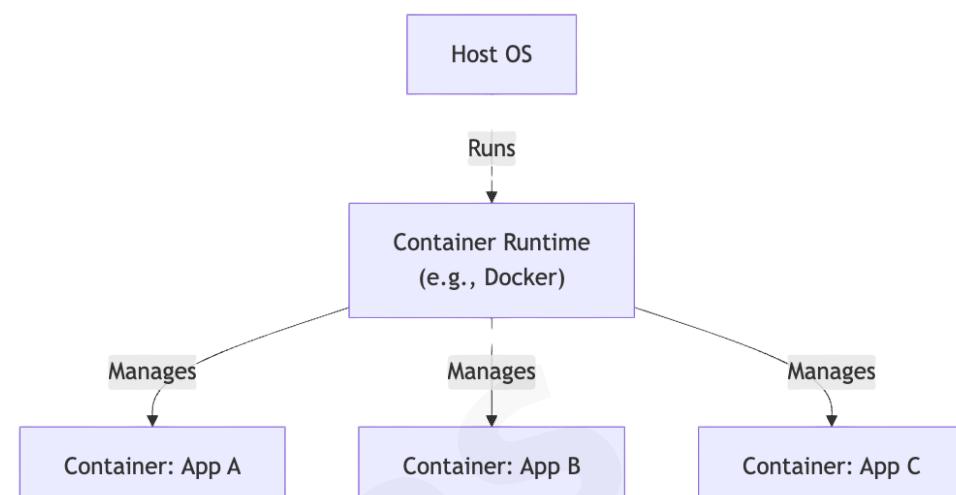
# Introduction

**Before Diving Deep into Kubernetes, UNDERSTAND this:**

## Containers:

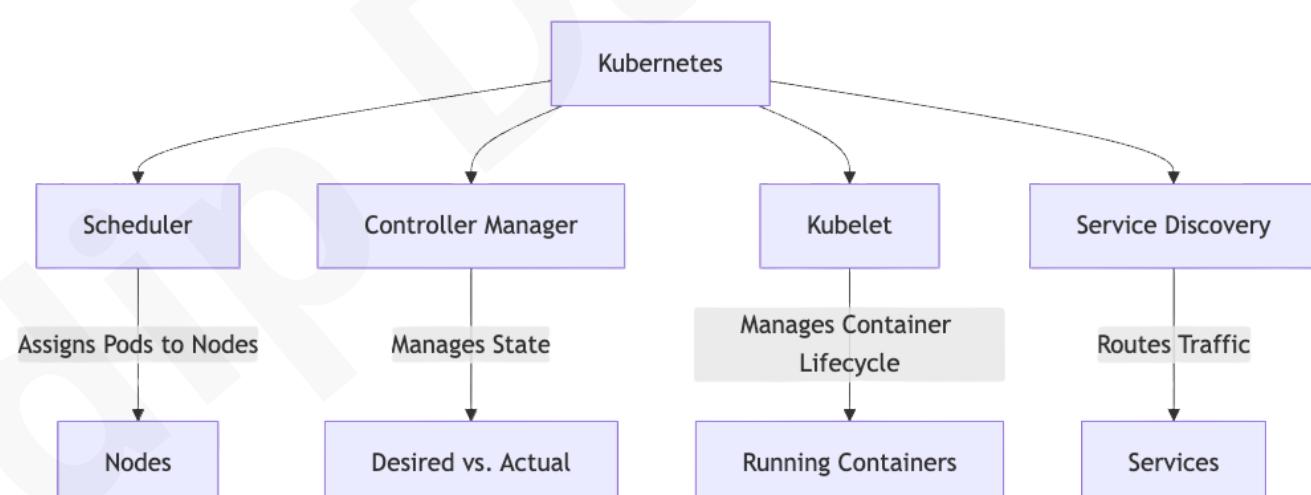
Containers are **lightweight**, **portable**, and **isolated** environments that **package** applications and their dependencies.

Kubernetes orchestrates multiple containers across different nodes in a cluster.



## Orchestration:

**Orchestration** means the **automation of container lifecycle** e.g. **deploy, scale, heal, network**

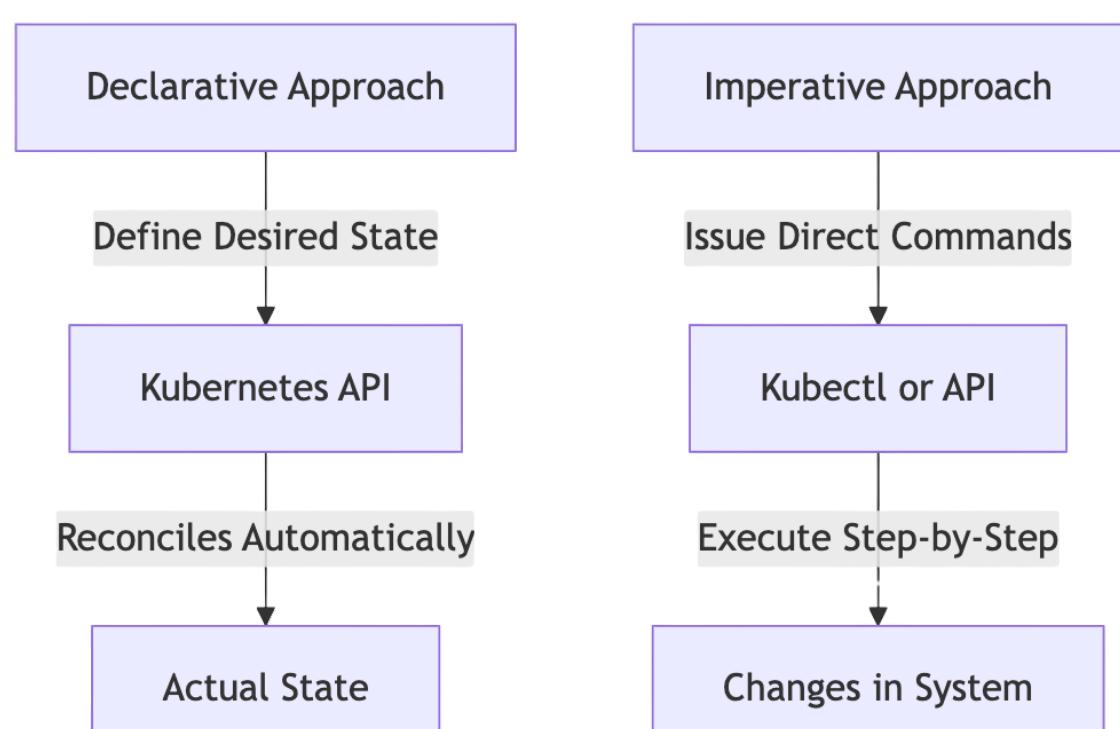


**Orchestration in Kubernetes** refers to the **automated deployment, scaling, and management** of containerized applications. Kubernetes ensures **high availability, self-healing, and efficient resource utilization**.

## Declarative vs. Imperative Approaches:

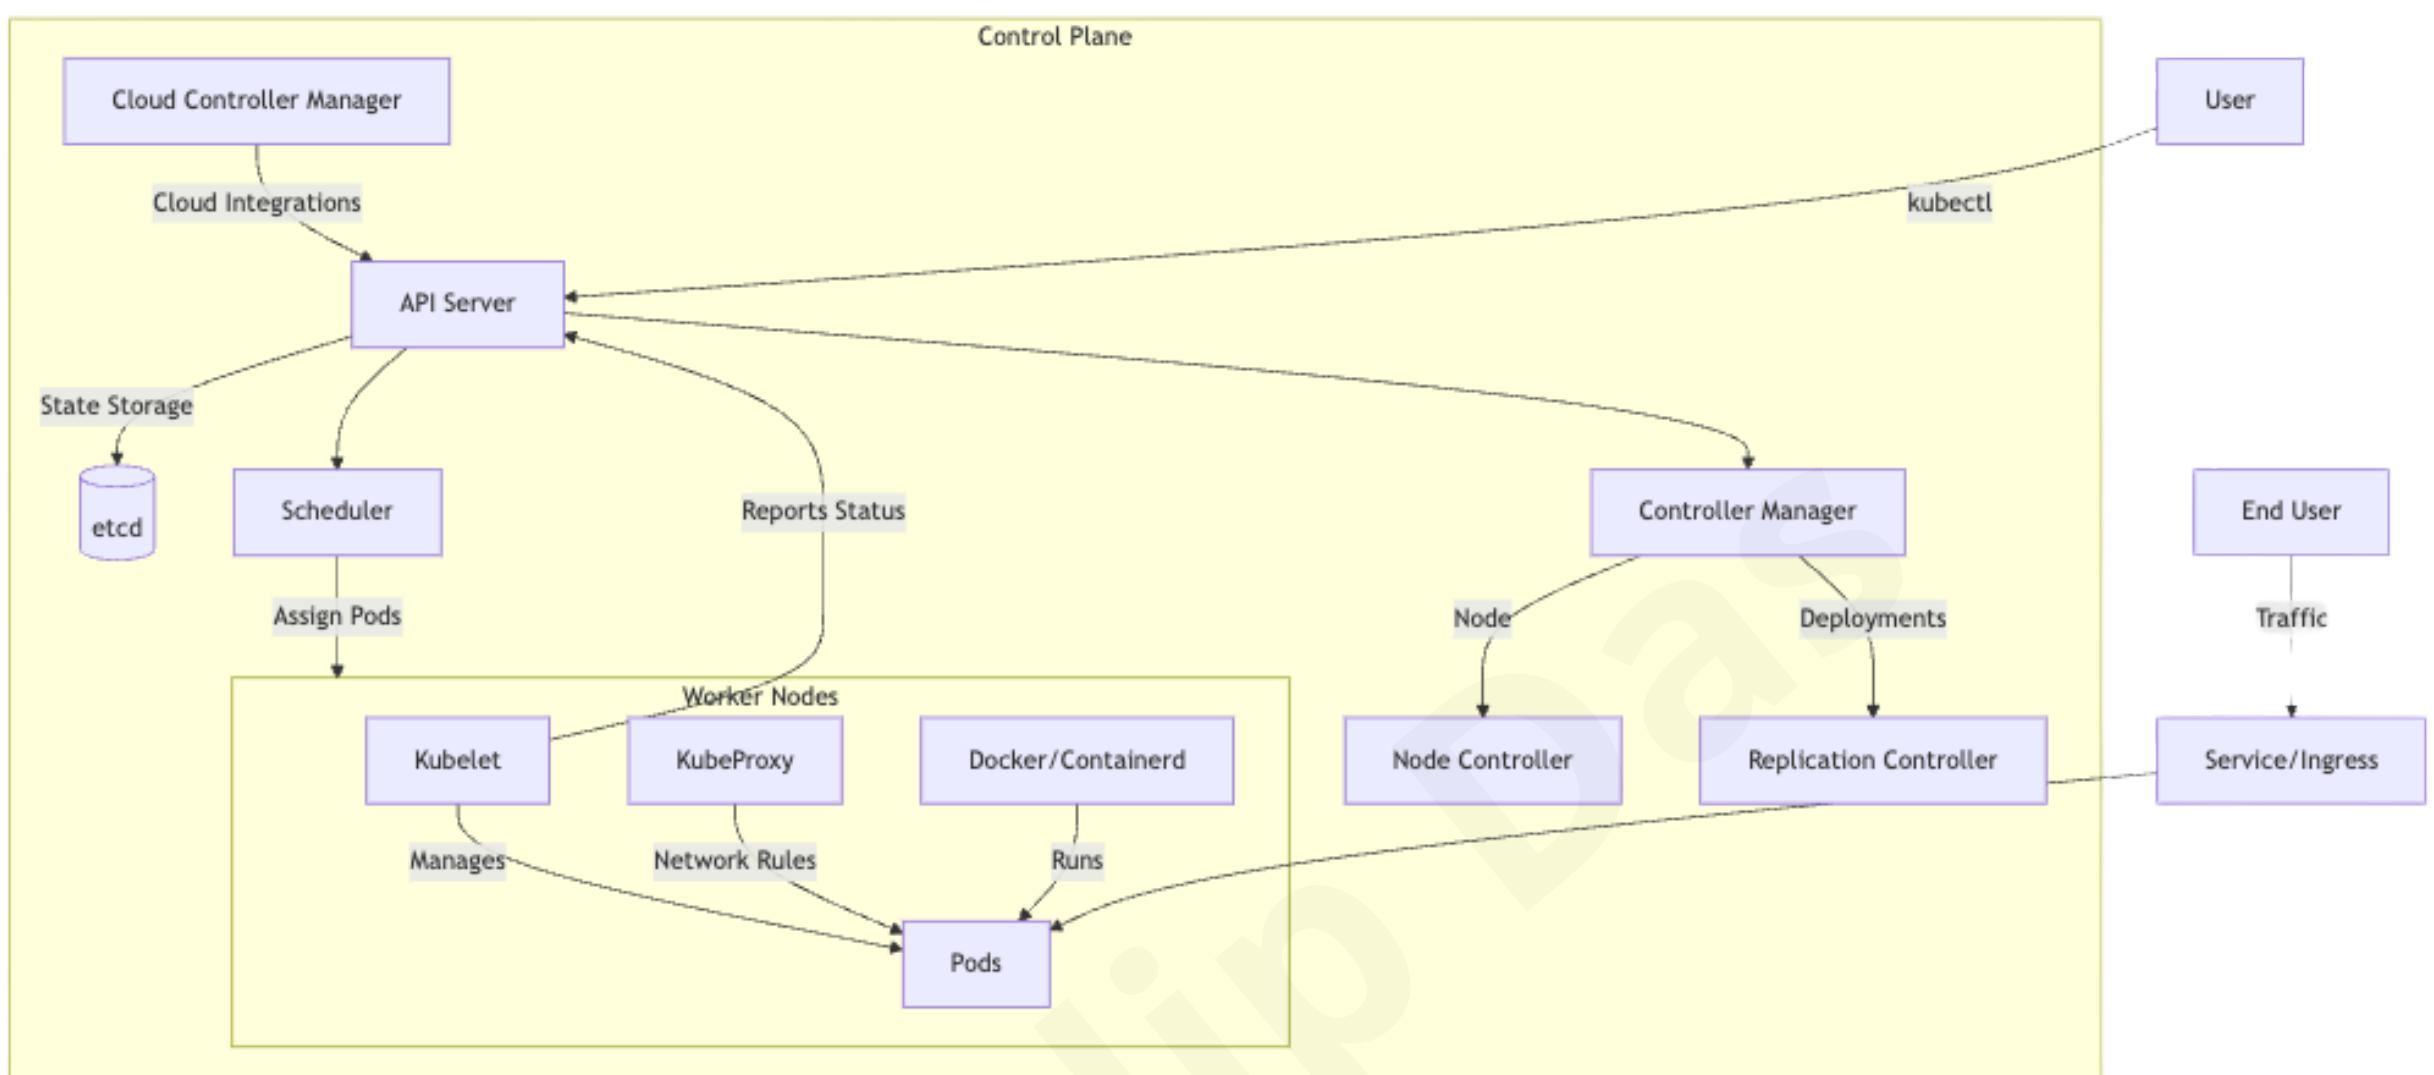
**Declarative:** Define desired state (YAML/JSON) and Kubernetes ensures it is maintained.

**Imperative:** Issue commands (e.g., kubectl create). Explicitly tell Kubernetes what to do at each step.



# Introduction

## Kubernetes Architecture Overview



### Control Plane (Master Nodes)

The Control Plane is responsible for managing and orchestrating the Kubernetes

- **API Server:** Frontend for cluster interactions (REST API).
- **etcd:** Distributed key-value store for cluster state.
- **Scheduler:** Assigns pods to nodes based on resources.
- **Controller Manager:** Runs core control loops (e.g., Node, ReplicaSet).
- **Cloud Controller Manager:** Manages cloud-specific integrations (e.g., load balancers).

### Worker Nodes

Worker nodes execute the actual application workloads by running containers inside Pods.

- **Kubelet:** Agent managing pods and reporting node status.
- **kube-proxy:** Maintains network rules for pod communication.
- **Container Runtime:** Runs containers (e.g., Docker, Containerd).

### Kubernetes Networking & Load Balancing / Add-ons

Networking in Kubernetes ensures seamless Pod-to-Pod communication, service discovery, and traffic routing.

- **Ingress** - Routes external traffic to services.
- **Service** - Provides stable network identity for Pods.
- **CNI (Container Network Interface)** - Ensures networking between Pods (Flannel, Calico).



# Introduction

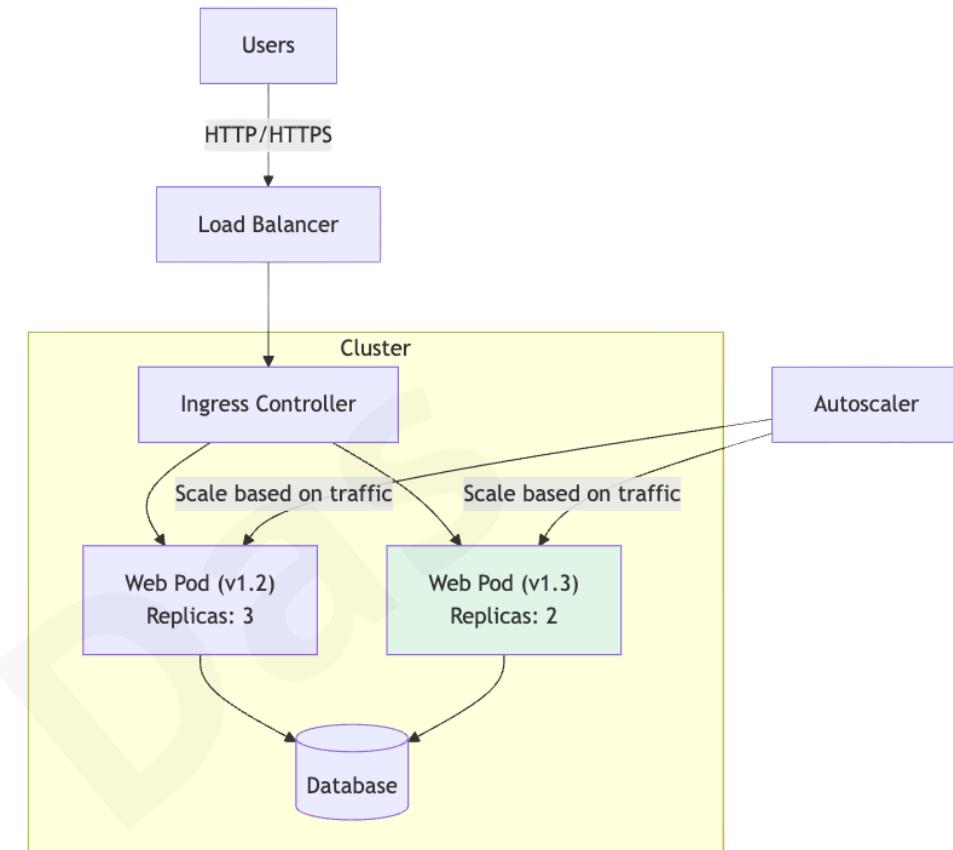
Now, let's discuss the Kubernetes usecases:

## Web Applications

Kubernetes efficiently manages scalable, resilient, and highly available web applications, automatically distributing traffic and handling failures.

### Key Features:

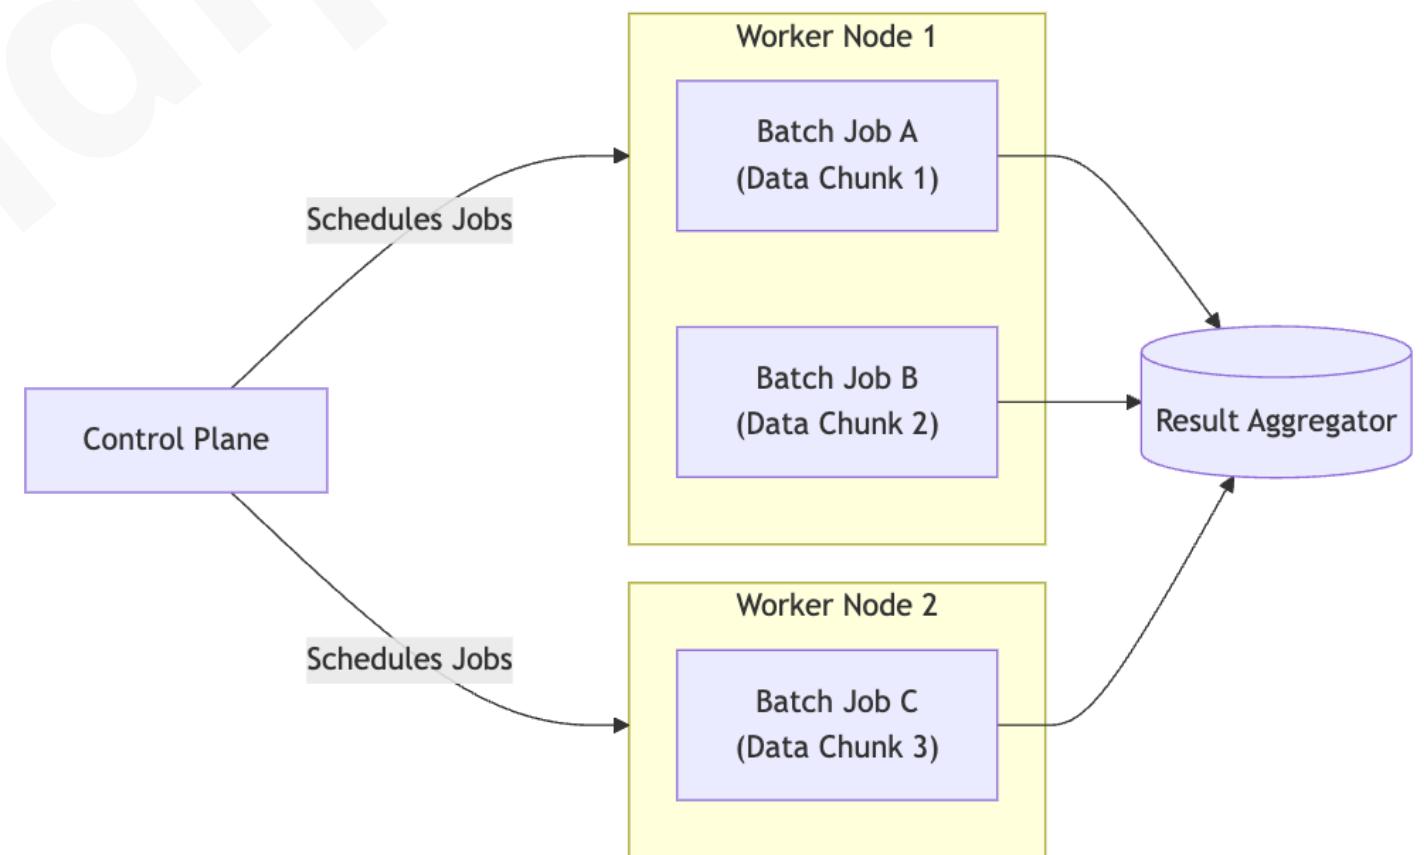
- **Load balancing** using Services and Ingress
- **Auto-scaling** with Horizontal Pod Autoscaler (HPA)
- **Database** connection with **Persistent Storage**



## Batch Processing

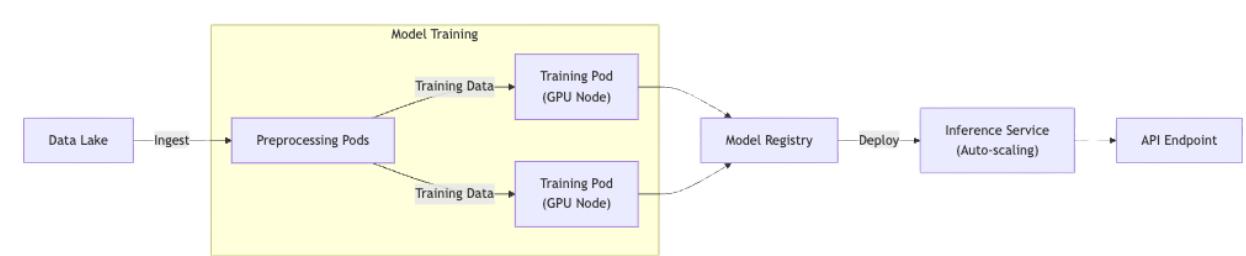
It's all about job **scheduling**, **parallel tasks**, and **resource management**

Kubernetes is ideal for batch processing jobs, handling periodic tasks such as data processing, video encoding, and report generation.



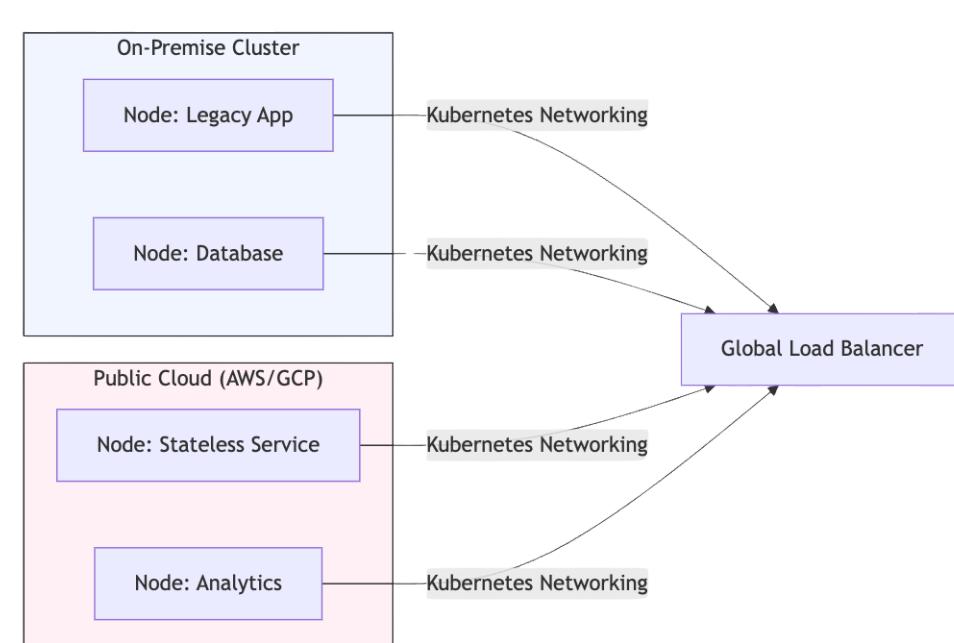
## Machine Learning

Kubernetes supports ML workloads by managing training jobs, inference services, and distributed computing using tools like Kubeflow.



## Hybrid Clouds

Kubernetes enables hybrid cloud architectures, allowing workloads to run across on-premises and cloud environments.



# Set-up Environment

## Local Clusters (For Development & Testing)

Ideal for learning, development, and testing applications locally before deploying to production.

- **Minikube**: Runs a single-node Kubernetes cluster locally.
- **Kind (Kubernetes in Docker)**: Uses Docker containers to simulate a multi-node Kubernetes cluster.
- **Docker Desktop**: Comes with a built-in Kubernetes cluster for quick local testing

## Production-Ready Self-Managed Clusters

For organizations that require full control over their Kubernetes infrastructure.

- **kubeadm**: Simplifies manual Kubernetes cluster setup.
- **kops**: Automates Kubernetes deployment on AWS.
- **Rancher**: Provides a user-friendly UI to manage Kubernetes clusters.

## Managed Kubernetes Services (Cloud Providers)

Best for scalable, secure, and automated production workloads without operational overhead.

- **Amazon EKS (Elastic Kubernetes Service)** - AWS-managed Kubernetes.
- **Google GKE (Google Kubernetes Engine)** - Google Cloud's Kubernetes service.
- **Azure AKS (Azure Kubernetes Service)** - Microsoft Azure's Kubernetes service.

### Cost Considerations:

- **EKS**: Control plane cost + EC2/node groups.
- **GKE**: Free control plane, charges for nodes + autopilot premium.
- **AKS**: Free control plane, pay only for nodes/services.

## Once again: When to Use What?

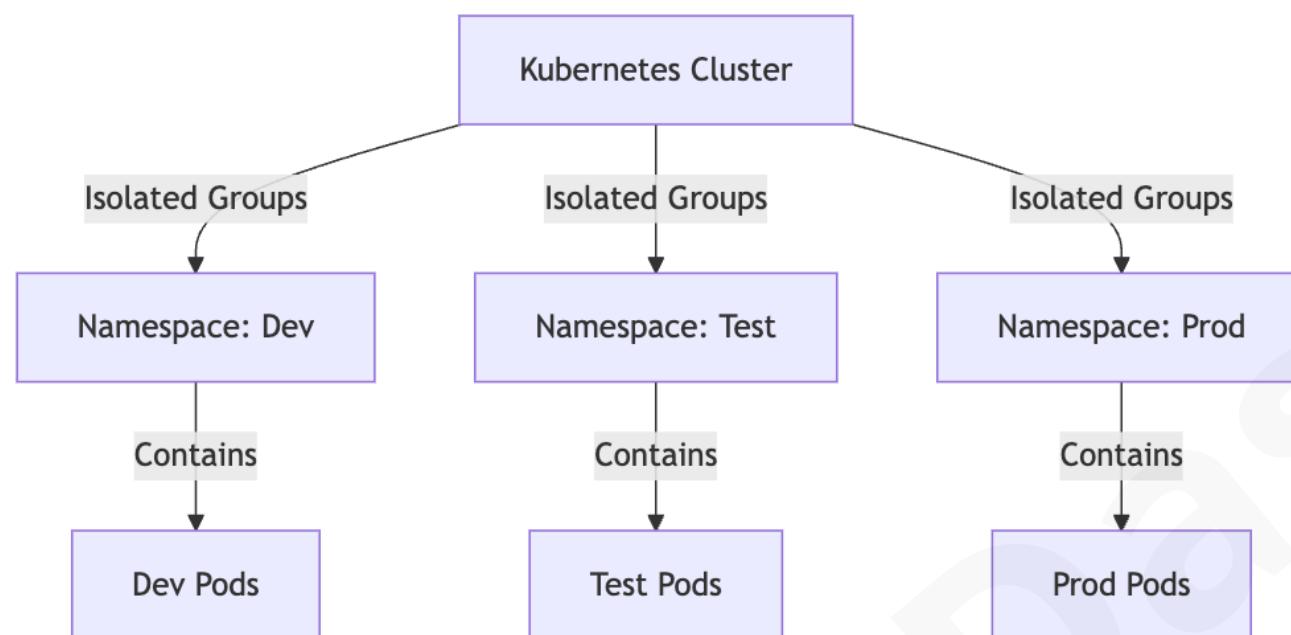
- **Local**: Learning, prototyping (minikube/kind).
- **Production Self-Managed**: Full control, compliance (kubeadm/kops).
- **Managed**: Speed, scalability, and reduced ops (EKS/GKE/AKS).



@Sandip Das

# Core Concepts in Action

Kubernetes provides several core concepts to efficiently manage containerized applications. Let's break down each concepts.



## Namespaces

**A Namespace is a virtual cluster within a Kubernetes cluster.**

### Key Use Cases

- Resource Isolation:** Has its own set of resources (Pods, Services, ConfigMaps, etc.), Prevent teams/apps from interfering (e.g., dev vs prod).
- Access Control:** RBAC rules scoped to namespaces.
- Resource Quotas:** Limit CPU/memory per namespace.

### Commands:

```
# List namespaces  
kubectl get namespaces  
  
# Create a namespace  
kubectl create namespace staging  
  
# Set default namespace for all commands  
kubectl config set-context --current --namespace=staging
```

### Common Default Namespaces

Namespace	Purpose
default	The default Namespace if none is specified.
kube-system	System-related Pods (DNS, controllers, etc.).
kube-public	Publicly readable resources (rarely used).
kube-node-lease	Stores heartbeats from nodes.
custom-namespace	User-defined Namespaces for applications.



@Sandip Das

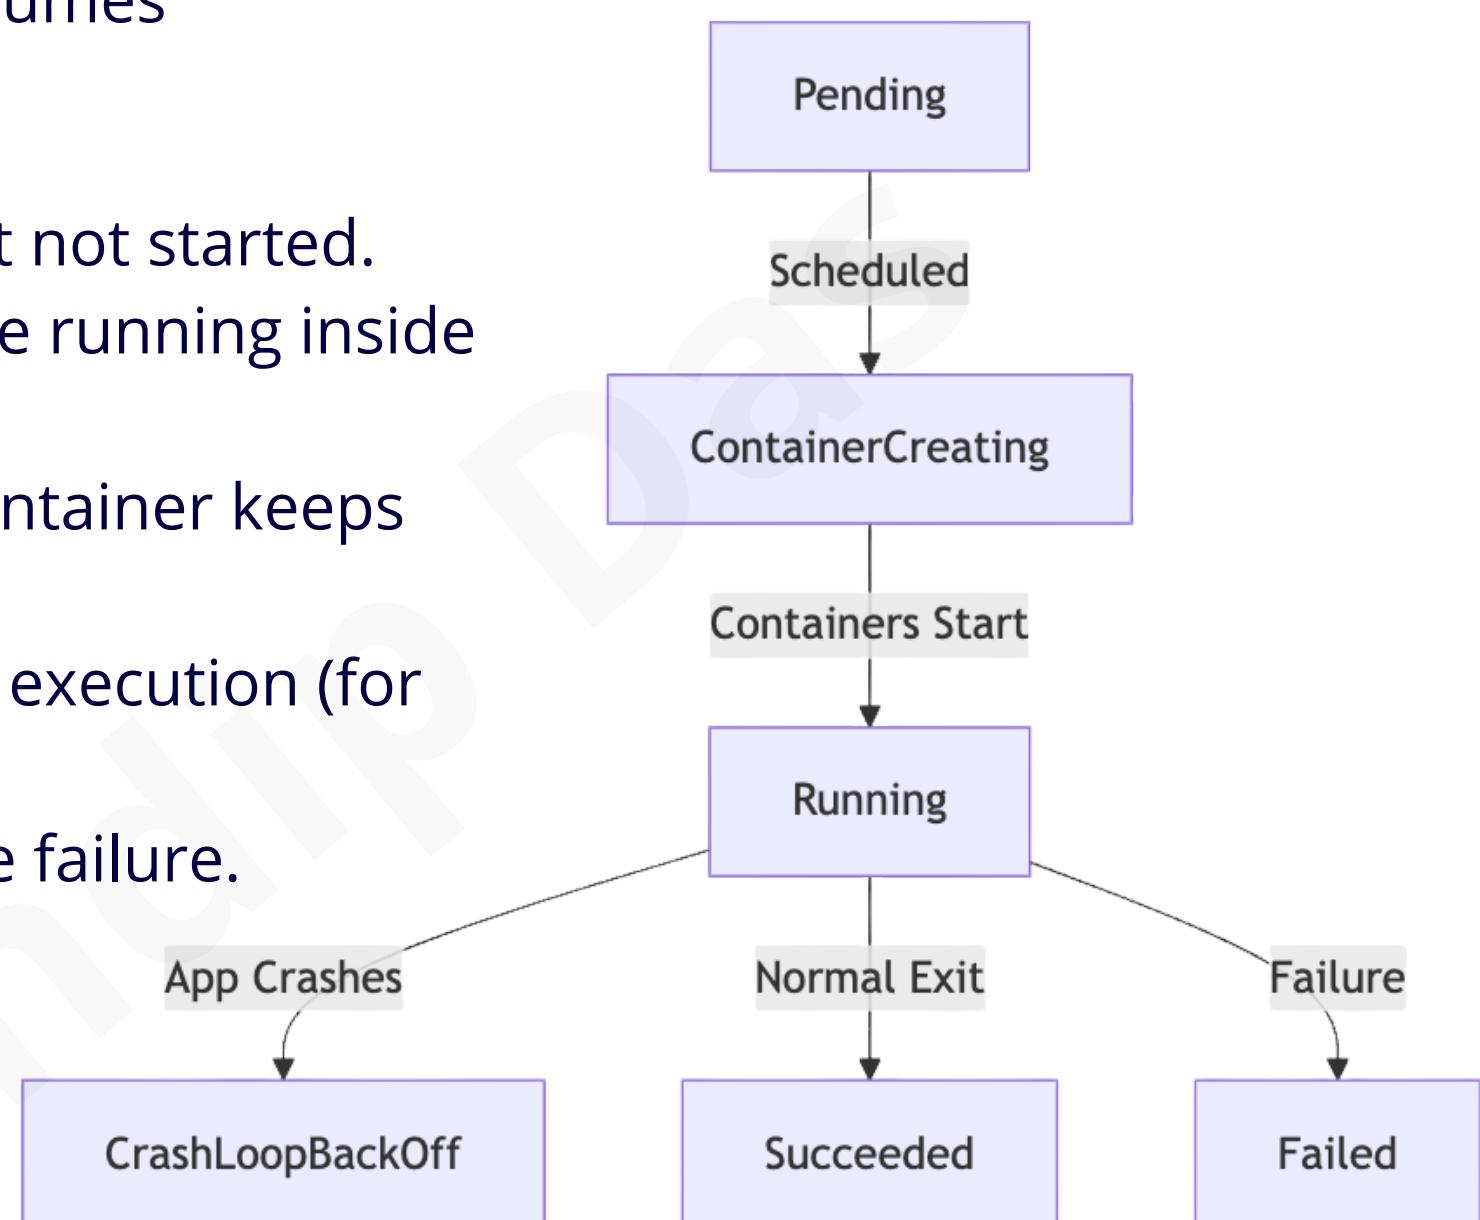
# Core Concepts in Action

## Pods: The Smallest Deployable Unit

A Pod encapsulates one or more containers, sharing the same network namespace and storage volumes

### Pod Lifecycle:

- **Pending:** Scheduled but not started.
- **Running:** Containers are running inside the Pod.
- **CrashLoopBackOff:** Container keeps failing and restarting.
- **Succeeded:** Completed execution (for jobs).
- **Failed:** Non-recoverable failure.

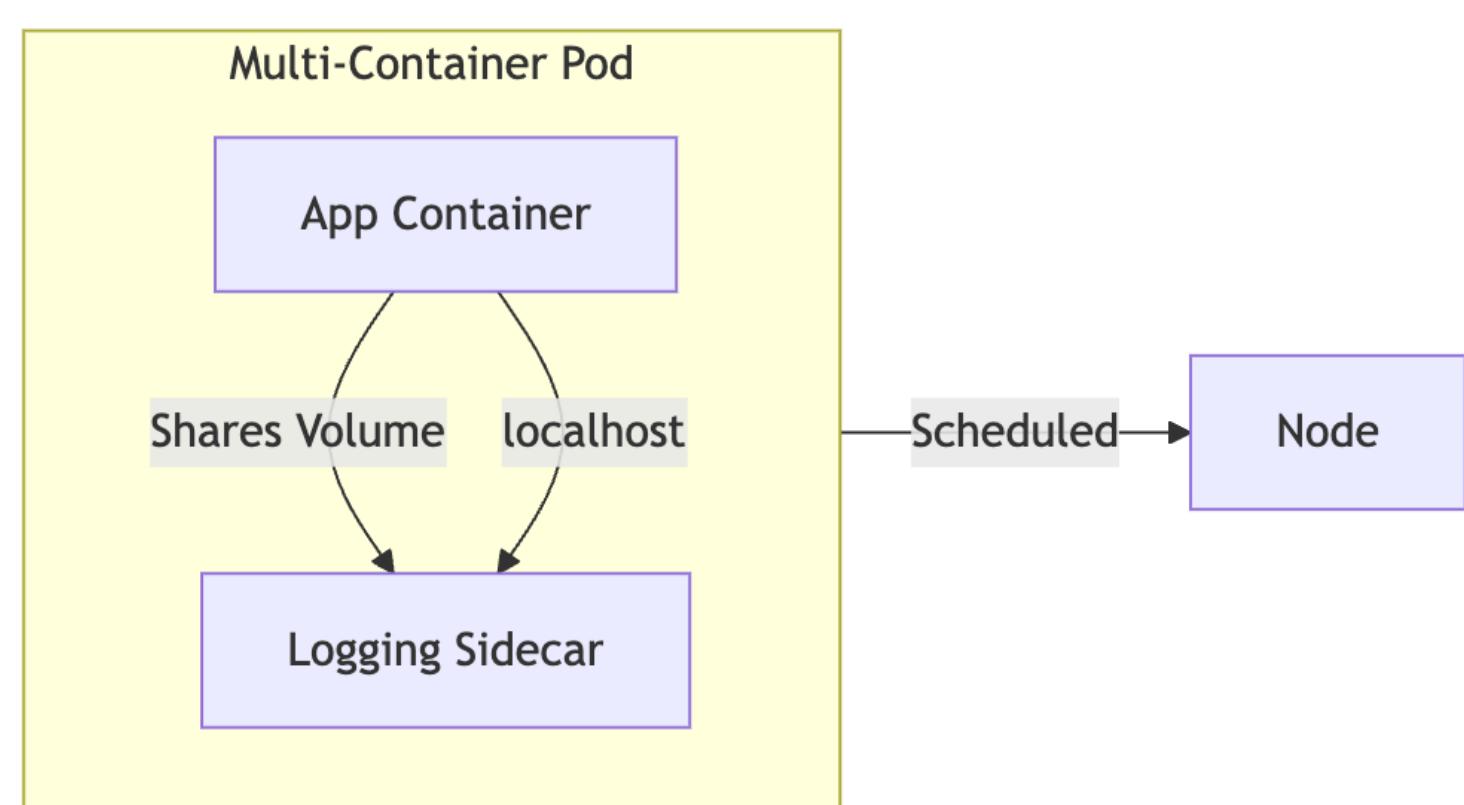


## Multi-Container Pods & Sidecar Pattern

A Pod can have multiple containers that communicate using localhost.

**The Sidecar pattern extends the primary container's functionality (e.g., logging, proxying). e.g.:**

```
apiVersion: v1
kind: Pod
metadata:
  name: sidecar-pod
spec:
  containers:
    - name: app-container
      image: my-app:latest
    - name: logging-sidecar
      image: fluentd
```



- The app container runs the main application.
- The sidecar container handles logging or metrics collection.



@Sandip Das

# Core Concepts in Action

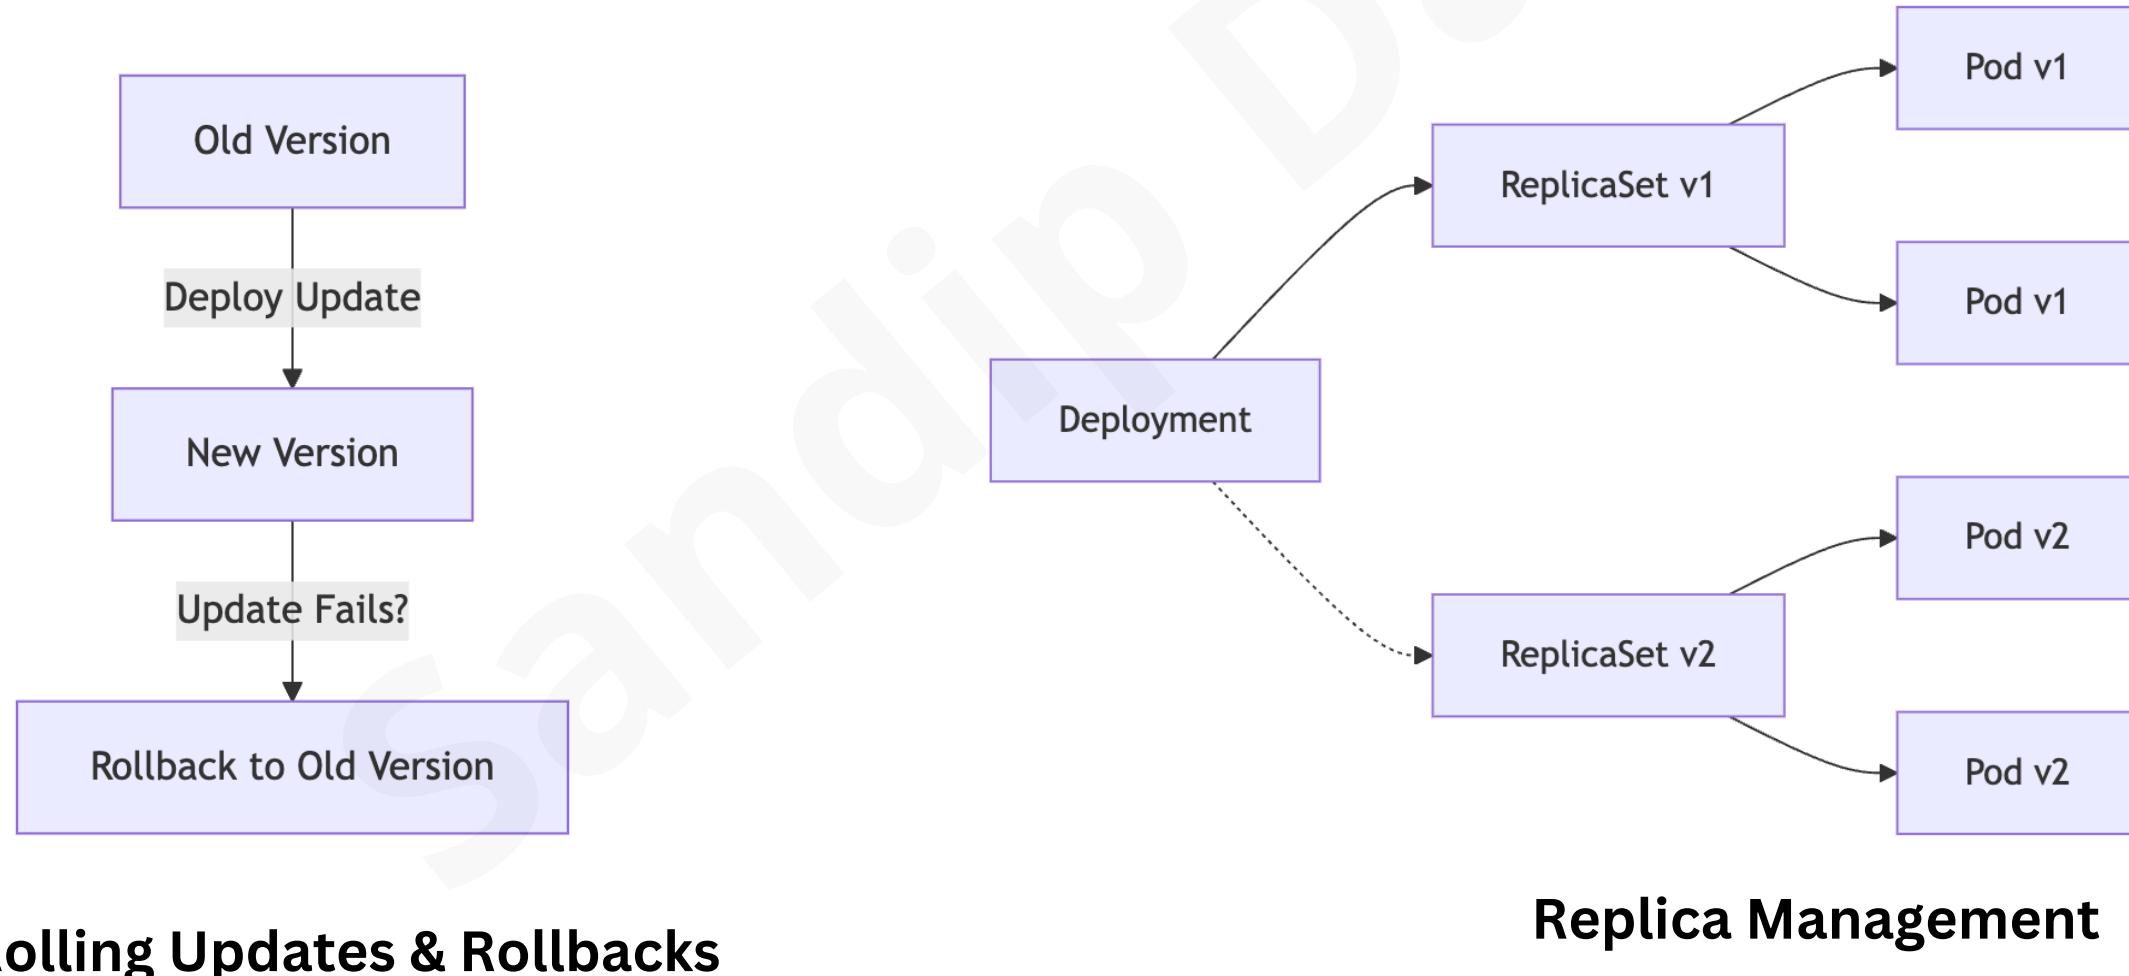
## Deployments: Managing Stateless Workloads

Deployments manage **rolling updates**, **rollbacks**, and **replica scaling**.

**Rolling Updates:** Gradual replacement of old pods with new ones.

**Rollbacks:** Revert to a previous ReplicaSet version.

**Replica Management:** A Deployment ensures the desired number of replicas is running.



### Example: Via :Imperative commands

```
kubectl create deployment nginx --image=nginx:1.20
```

```
kubectl set image deployment/nginx nginx=nginx:1.21 # Rolling update
```

```
kubectl rollout undo deployment/nginx # Rollback
```

nginx\_deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
```

### Example: Via :Declarative

We can do via Kubernetes manifest files , run:

```
kubectl apply -f nginx_deployment.yaml
```



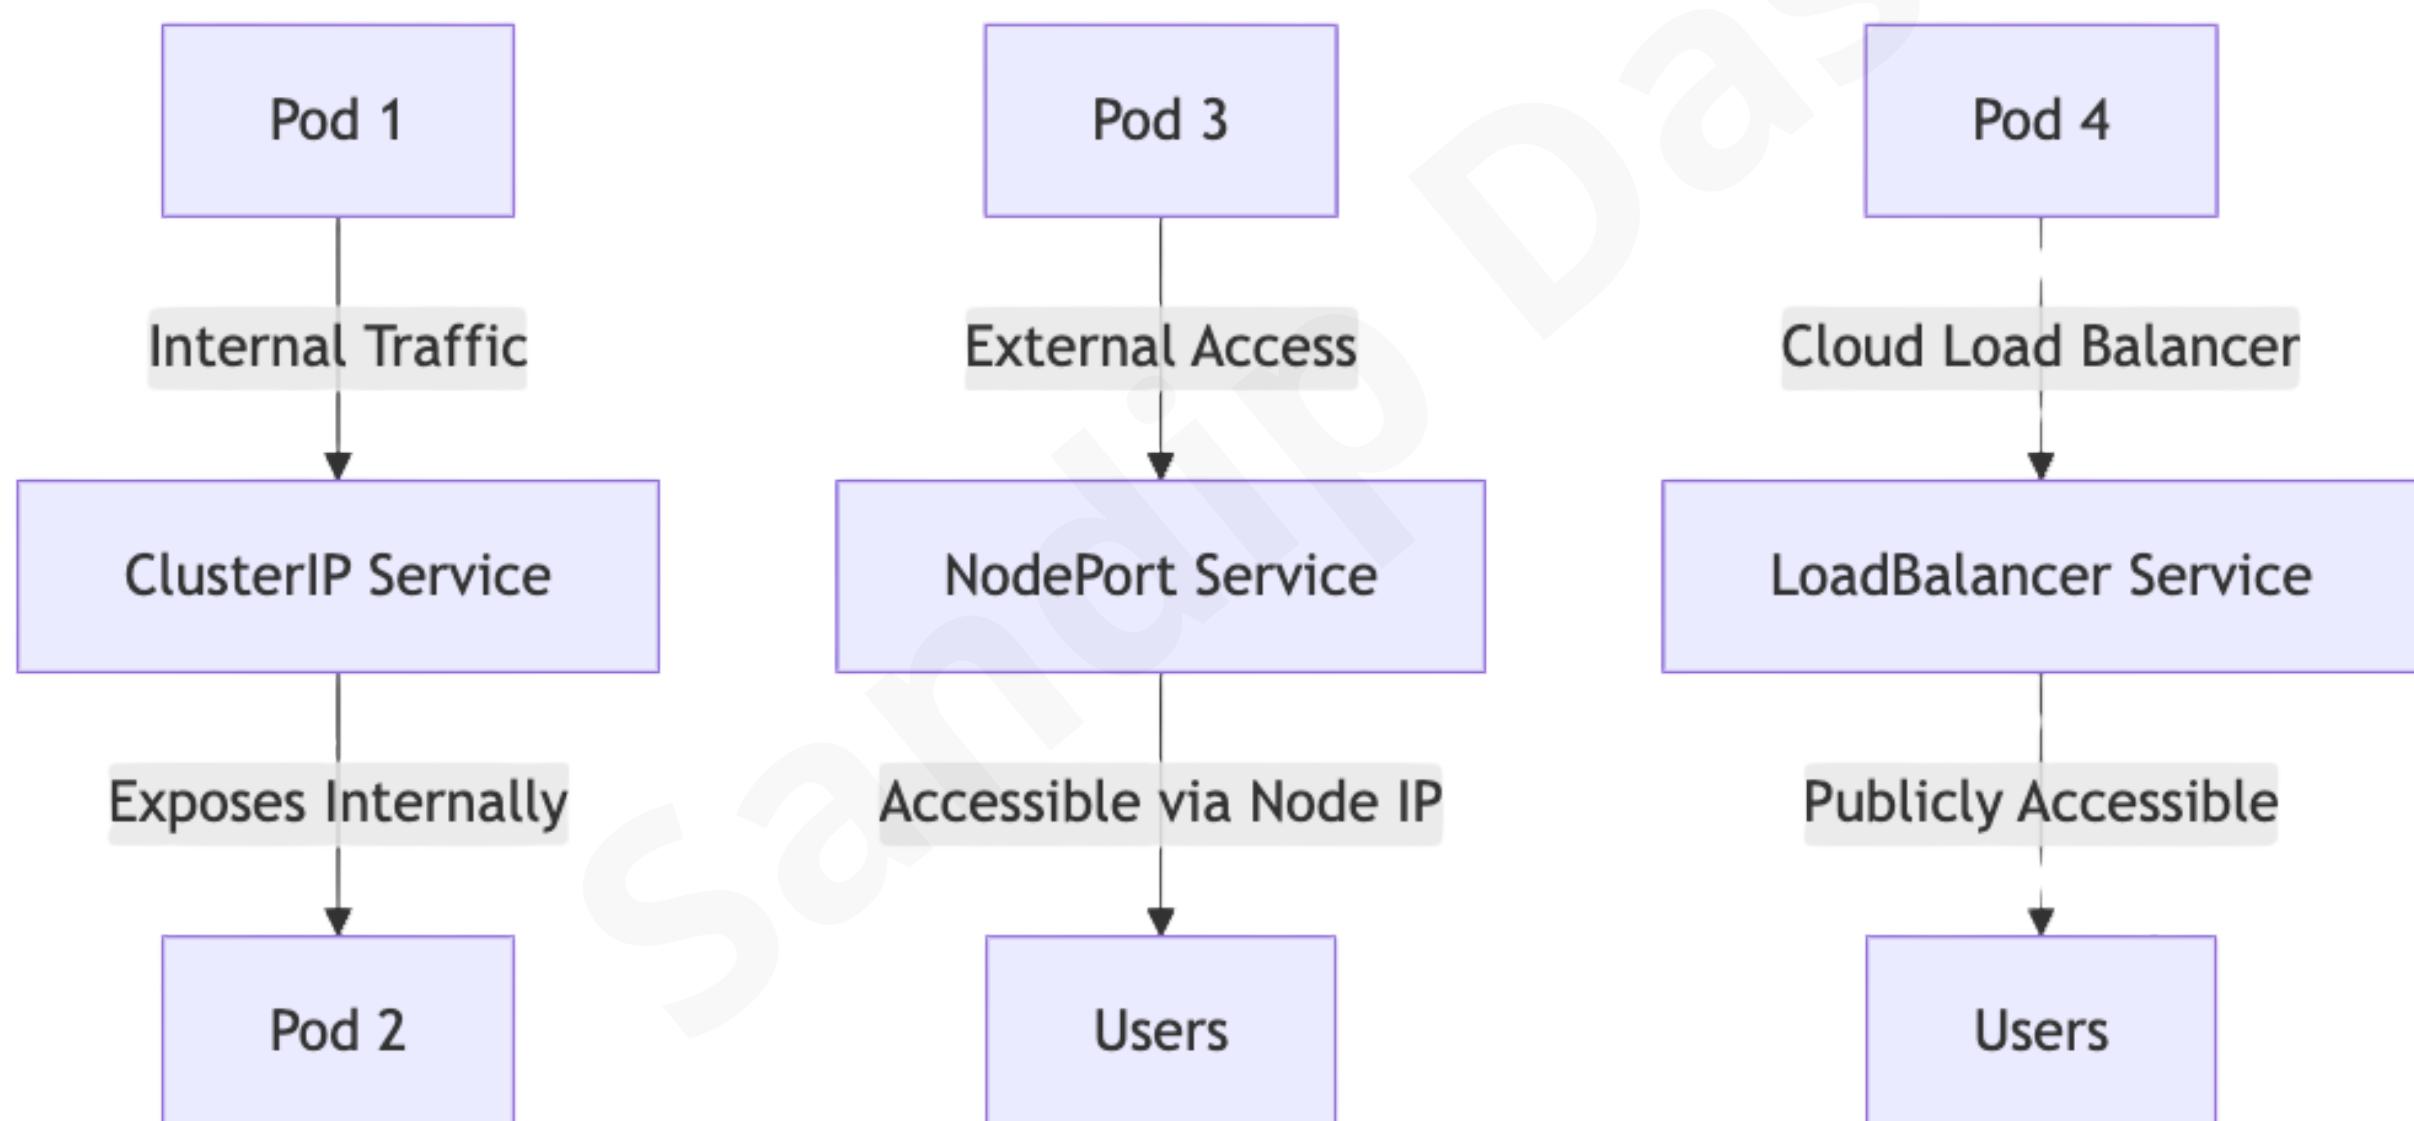
@Sandip Das

# Core Concepts in Action

## Services

A Service enables communication between Pods and external users.

- **ClusterIP:** Internal IP for intra-cluster communication.
- **NodePort:** Exposes a specific port on each node.
- **LoadBalancer:** Uses a cloud provider's load balancer.



```
apiVersion: v1
kind: Service
metadata:
  name: web-service
spec:
  type: ClusterIP
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  type: NodePort
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
      nodePort: 30080
```

```
apiVersion: v1
kind: Service
metadata:
  name: web-service
spec:
  type: LoadBalancer
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```



# Core Concepts in Action

## ConfigMaps & Secrets

**ConfigMaps** store non-sensitive configuration data, while **Secrets** handle sensitive data.

**ConfigMap: Non-sensitive configs** (e.g., environment variables).

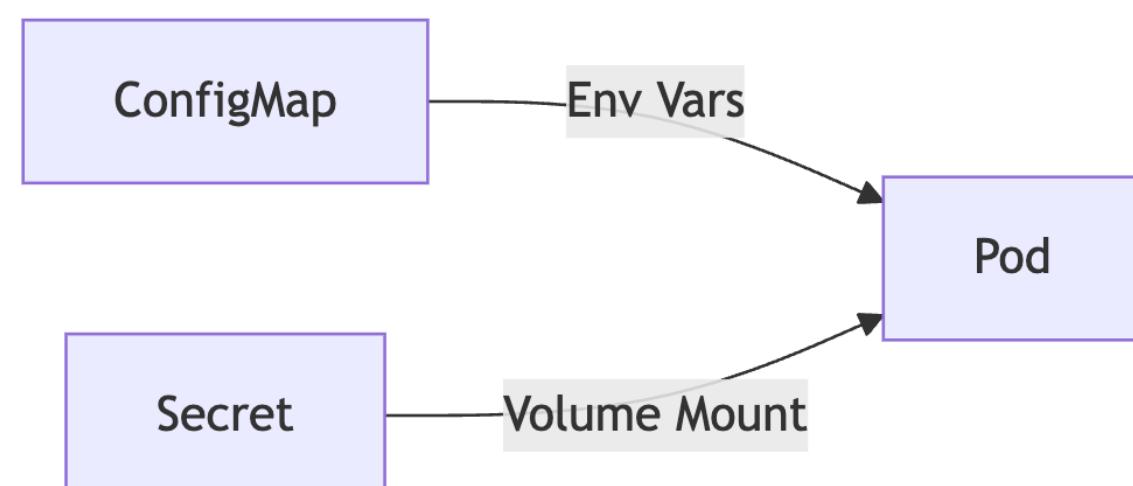
**Secret: Sensitive data** (base64-encoded).

### Example: ConfigMap (Environment Variables)

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  DATABASE_URL: "mysql://db:3306"
```

### Mount the ConfigMap as an environment variable:

```
env:
- name: DATABASE_URL
valueFrom:
  configMapKeyRef:
    name: app-config
    key: DATABASE_URL
```



### Example: Secret (Base64 Encoded Data)

```
apiVersion: v1
kind: Secret
metadata:
  name: mysql-secret
type: Opaque
data:
  password: cGFzc3dvcmQ= # Base64 encoded "password"
```

### Mount the Secret as an environment variable:

```
env:
- name: MYSQL_PASSWORD
valueFrom:
  secretKeyRef:
    name: mysql-secret
    key: password
```



@Sandip Das

# Advanced Workloads and Patterns

Kubernetes provides advanced workload controllers and patterns for managing stateful applications, node-wide services, batch processing, and custom automation. Let's dive into:

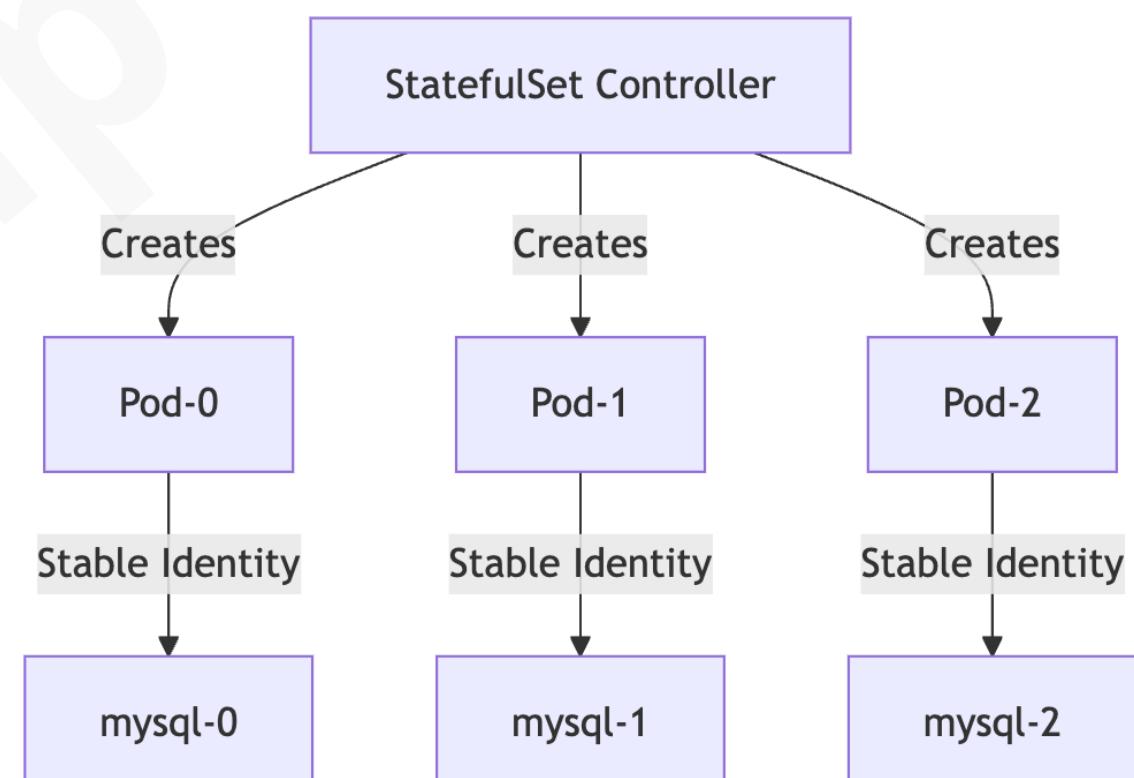
- **StatefulSets** (for databases)
- **DaemonSets** (for node monitoring)
- **Jobs & CronJobs** (for batch processing)
- **Operators & CRDs** (for automation)

## StatefulSets

Unlike Deployments, **StatefulSets** maintain sticky identities for Pods, making them ideal for databases like Cassandra, MySQL, MongoDB.

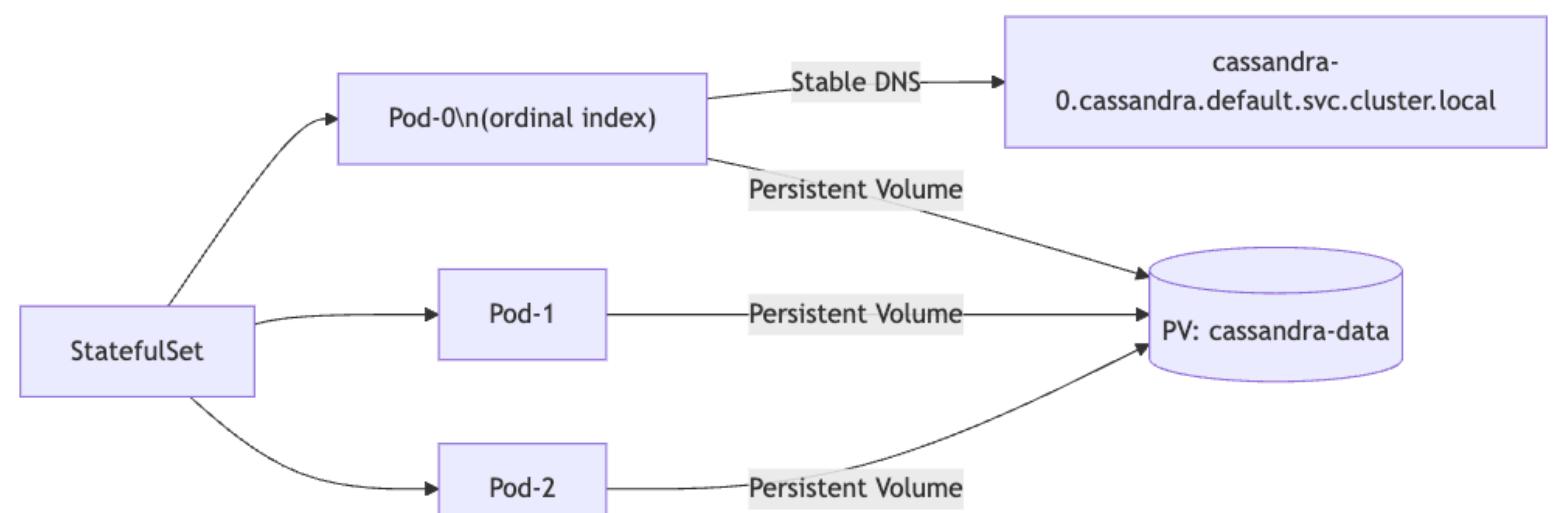
Key Features:

- Unique, stable network identifiers (pod-0, pod-1, pod-2...).
- Persistent storage maintained across pod restarts.
- Ordered pod creation, scaling, and deletion.



## Example: StatefulSet for Cassandra

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: cassandra
spec:
  serviceName: "cassandra"
  replicas: 3
  selector:
    matchLabels:
      app: cassandra
  template:
    metadata:
      labels:
        app: cassandra
    spec:
      containers:
        - name: cassandra
          image: cassandra:latest
          ports:
            - containerPort: 9042
```



[Click here for more advance example](#)



@Sandip Das

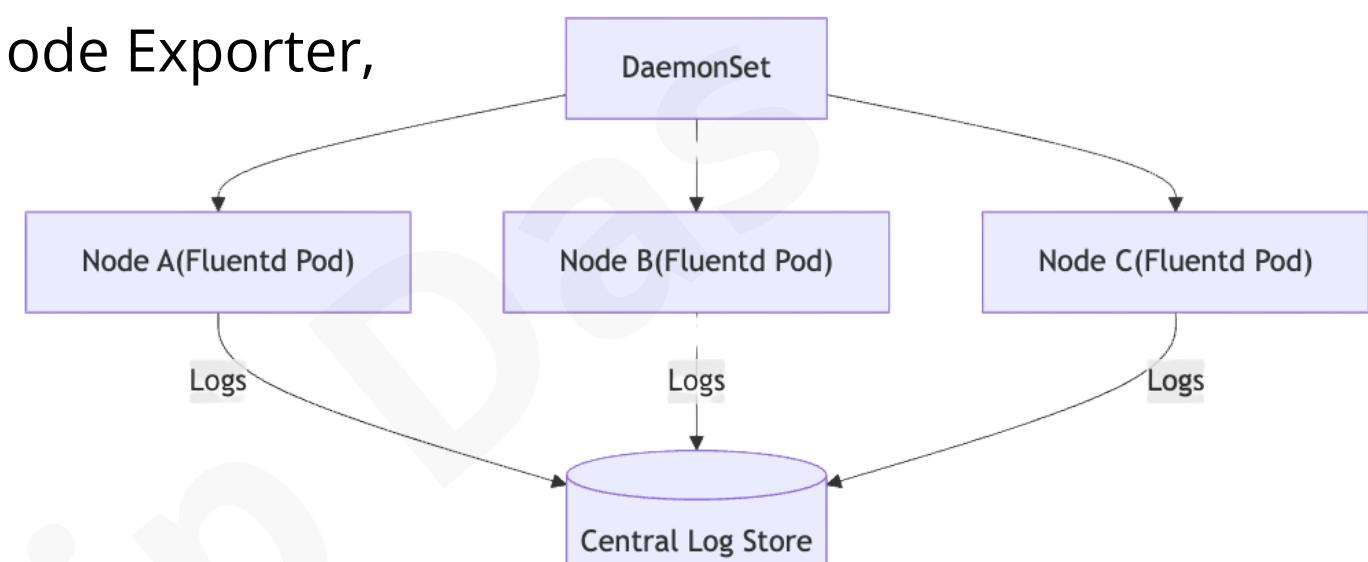
# Advanced Workloads and Patterns

## DaemonSets

A DaemonSet ensures that one copy of a Pod runs on every node in the cluster. Useful for node monitoring, logging, and security agents.

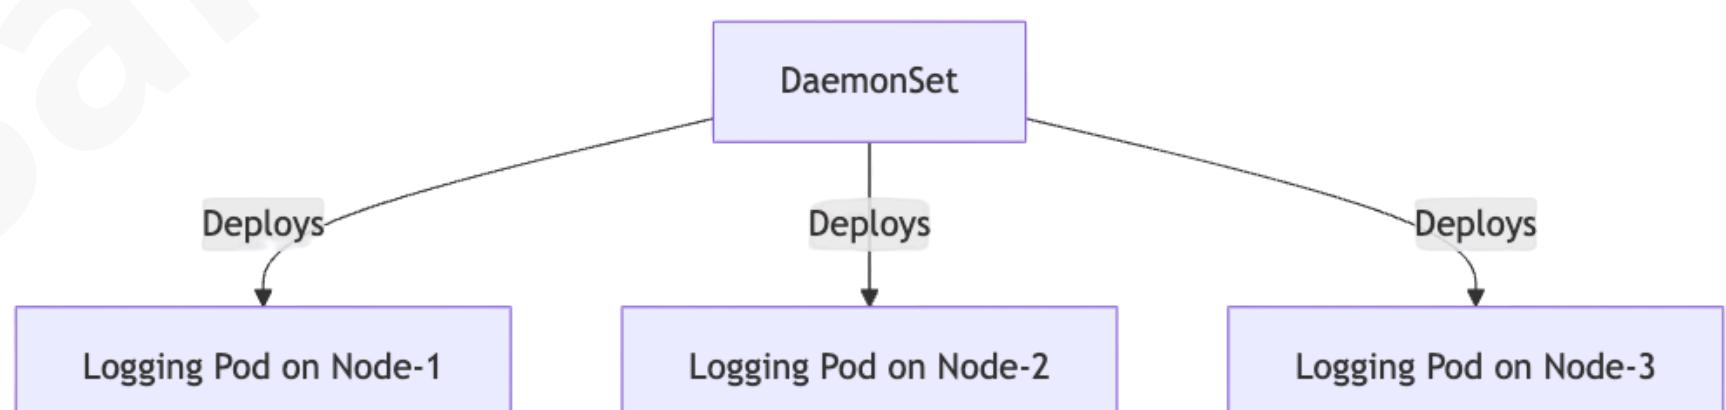
### Common Use Cases:

- Log collectors (Fluentd, Filebeat).
- Node monitoring (Prometheus Node Exporter, Datadog).
- Security agents (Falco, Sysdig).



### Example: Fluentd DaemonSet

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd
spec:
  selector:
    matchLabels:
      name: fluentd
  template:
    metadata:
      labels:
        name: fluentd
    spec:
      containers:
        - name: fluentd
          image: fluent/fluentd-kubernetes-daemonset
          volumeMounts:
            - name: varlog
              mountPath: /var/log
      volumes:
        - name: varlog
          hostPath:
            path: /var/log
```



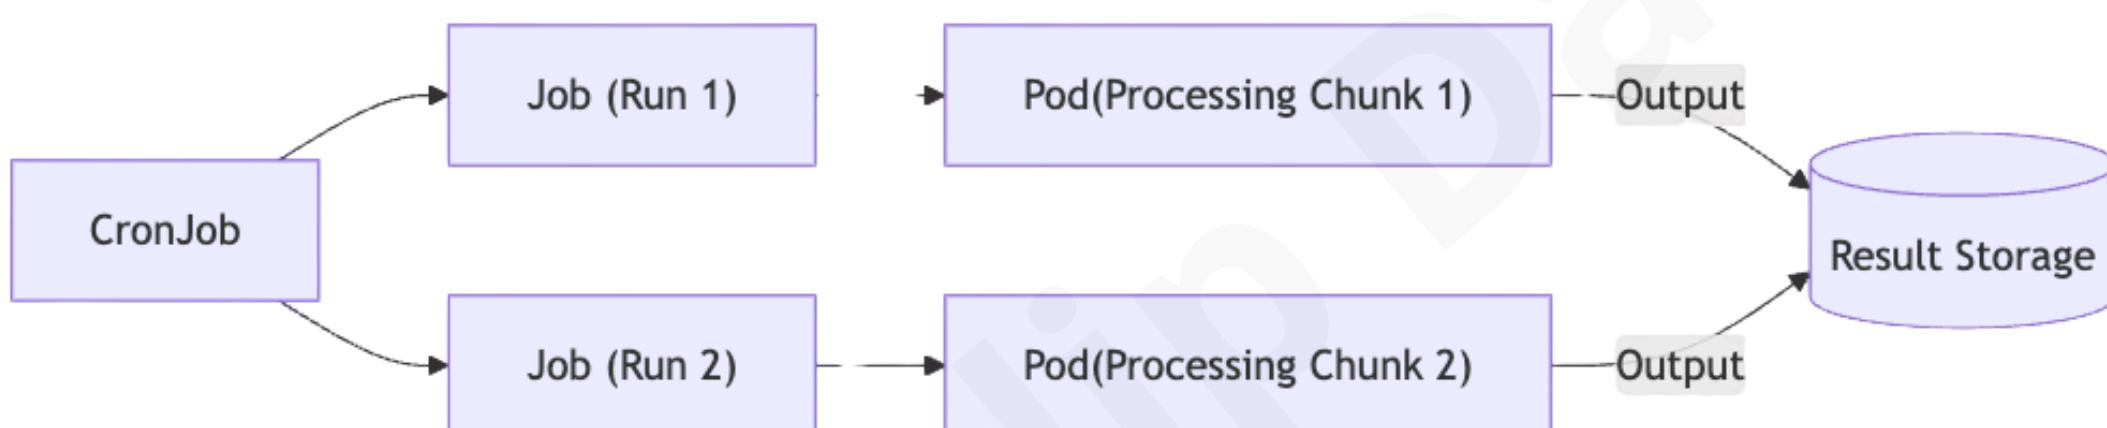
# Advanced Workloads and Patterns

## Jobs & CronJobs

A Job runs a task until completion, while a CronJob schedules recurring tasks. (e.g., data processing, backups).

### Key Features:

- ✓ Runs to completion (no restart on failure).
- ✓ Ideal for ETL jobs, report generation, backups.
- ✓ Supports parallelism and retry policies.



### Example (CronJob for daily cleanup):

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: daily-cleanup
spec:
  schedule: "0 0 * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: cleanup
              image: alpine
              command: ["/bin/sh", "-c", "rm -rf /tmp/old-files"]
        restartPolicy: OnFailure
```

### Example: Job (One-Time Execution):

```
apiVersion: batch/v1
kind: Job
metadata:
  name: data-processing-job
spec:
  completions: 3
  template:
    spec:
      containers:
        - name: data-processor
          image: python:3.8
          command: ["python", "-c", "print('Processing data...')"]
  restartPolicy: Never
```



@Sandip Das

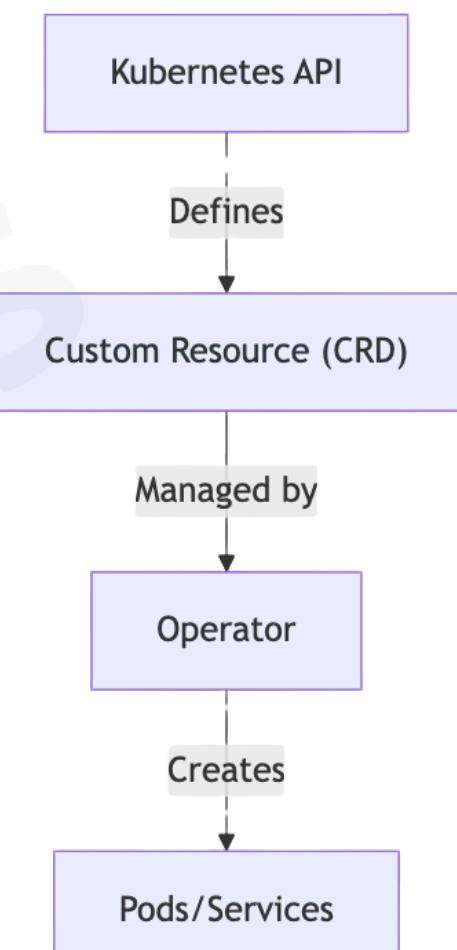
# Advanced Workloads and Patterns

## Operators & Custom Resource Definitions (CRDs)

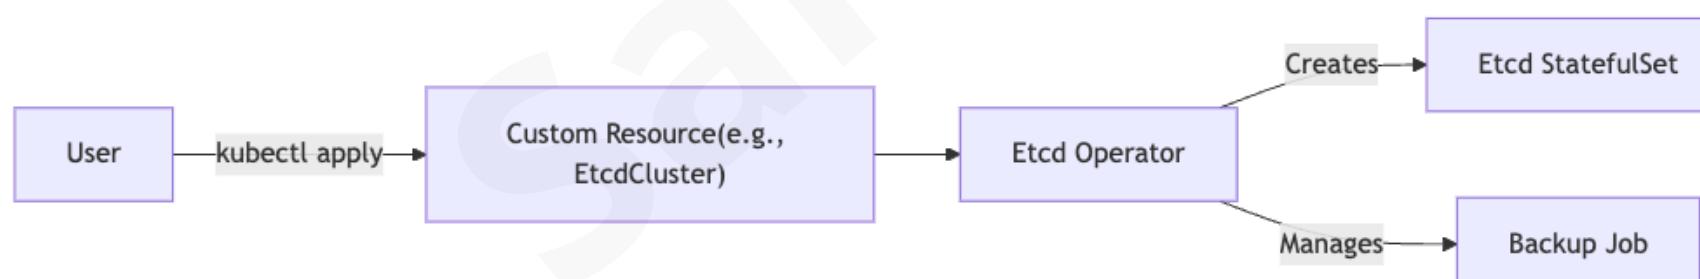
Operators automate complex application lifecycle management by defining custom Kubernetes controllers (e.g., databases, middleware).

### Key Features:

- Automates application deployment, scaling, backups, failover.
- Uses Custom Resource Definitions (CRDs) to extend Kubernetes API.
- Ideal for managing databases, ML workflows, monitoring tools.



### Example (Etcd Operator):



#### Step 1: Define a CRD:

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: etcdclusters.etcd.database.coreos.com
spec:
  group: etcd.database.coreos.com
  scope: Namespaced
  names:
    plural: etcdclusters
    singular: etcdcluster
    kind: EtcdCluster
  versions:
    - name: v1beta2
      served: true
      storage: true
```

#### Step 2: Deploy the Operator::

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: etcd-operator
spec:
  replicas: 1
  template:
    spec:
      containers:
        - name: etcd-operator
          image: quay.io/coreos/etcd-operator:v0.9.4
```

#### Step 3: Create a Custom Resource:

```
apiVersion: etcd.database.coreos.com/v1beta2
kind: EtcdCluster
metadata:
  name: example-etcd-cluster
spec:
  size: 3
  version: 3.4.9
```



# Networking Deep Dive

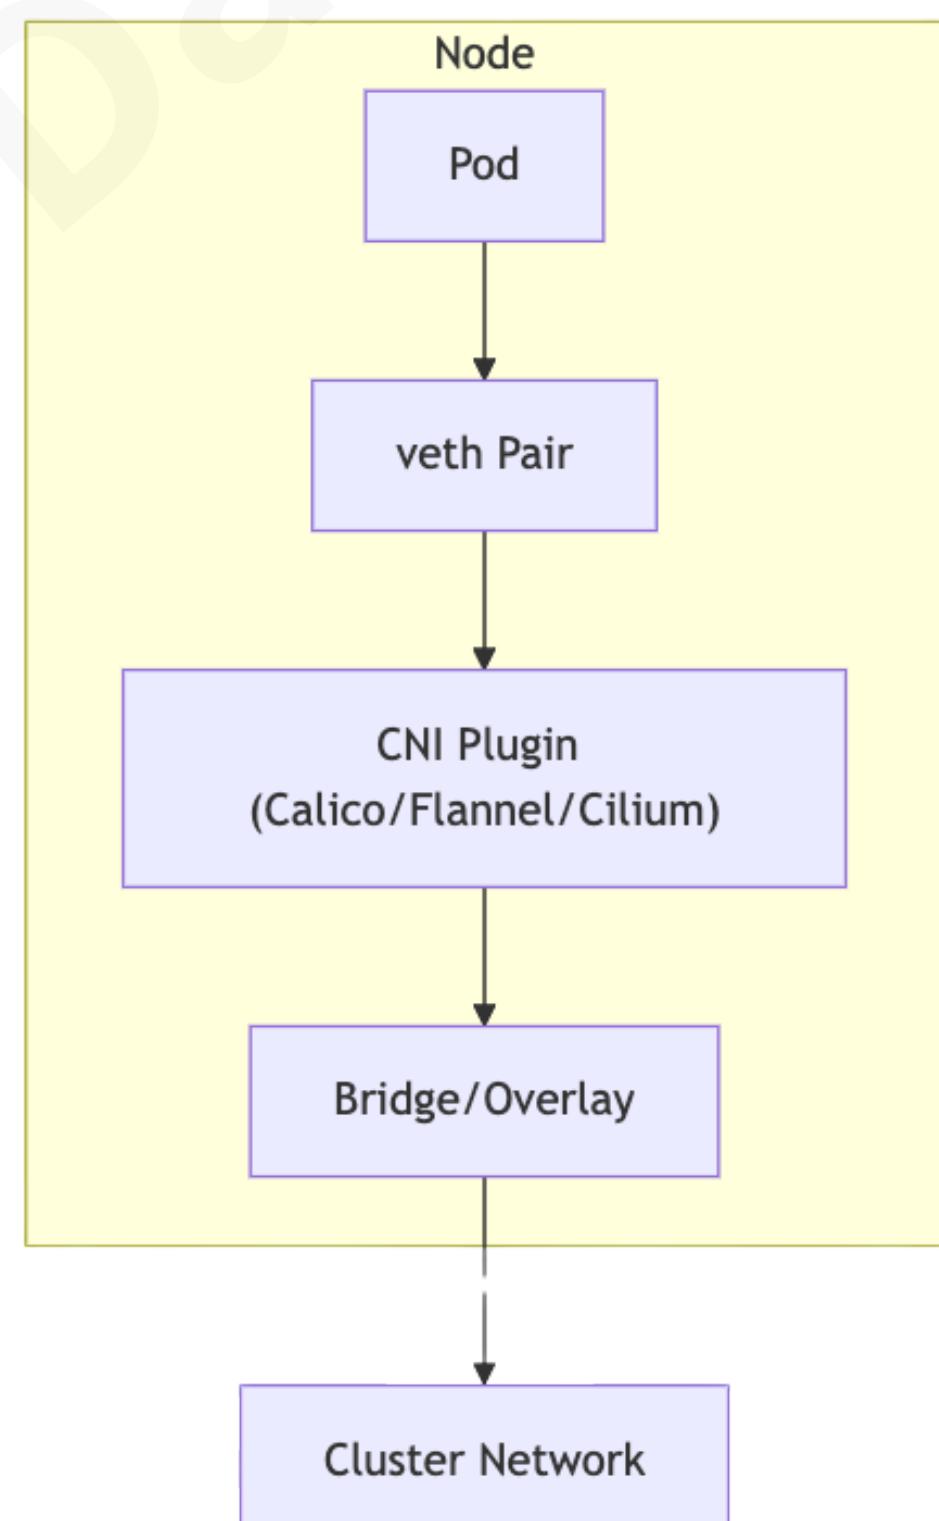
Kubernetes networking is a critical aspect of managing pod communication, service discovery, ingress traffic, and security. Let's dive into:

- CNI Plugins (Calico, Flannel, Cilium)
- Ingress Controllers (NGINX, Traefik)
- DNS & Service Discovery (CoreDNS, troubleshooting)
- Hands-on Lab (Expose a microservice via Ingress)

## CNI Plugins: Enabling Pod-to-Pod Networking

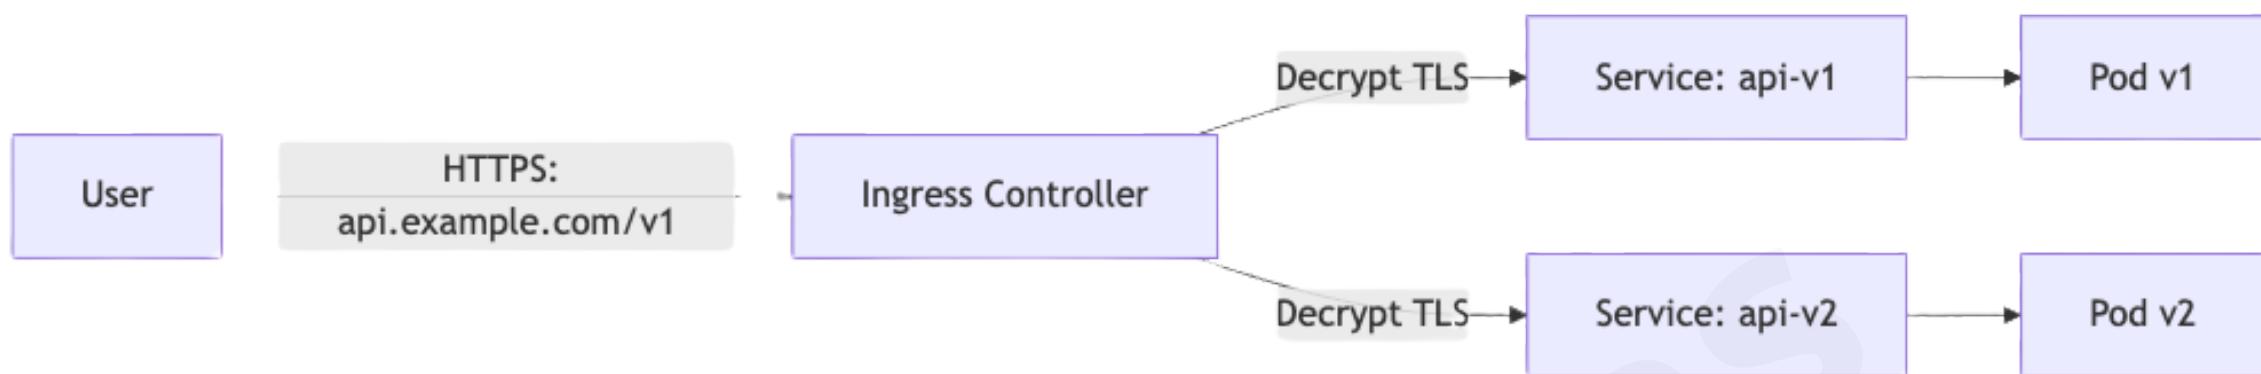
**Kubernetes relies on Container Network Interface (CNI) plugins to manage networking between Pods.**

Plugin	Key Features	Use Case
Calico	BGP routing, network policies, IPv6/IPv4 support.	Enterprise security/compliance.
Flannel	Simple overlay network (VXLAN or host-gw).	Basic networking for small clusters.
Cilium	eBPF-based networking, L7 policies, service mesh.	High-performance observability.



# Networking Deep Dive

# Ingress Controllers: Managing External Traffic



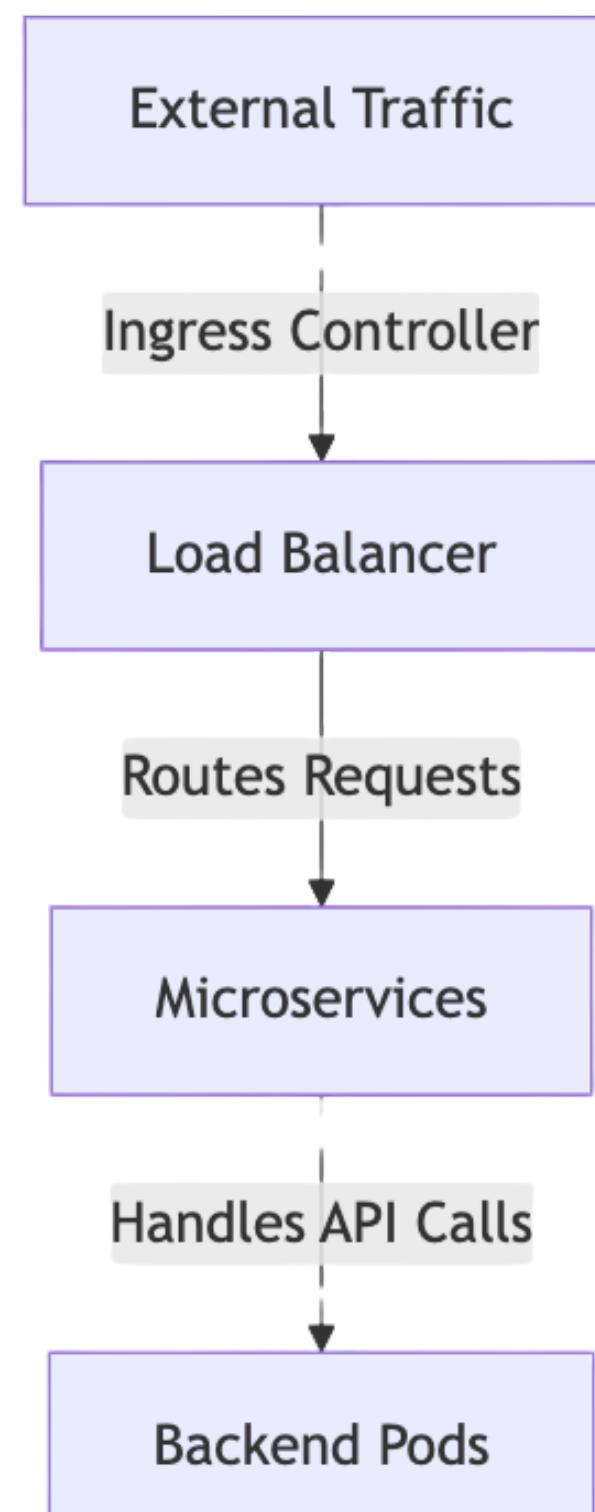
**Ingress routes HTTP/HTTPS traffic to services inside the cluster.**

# Popular Ingress Controllers

- **NGINX Ingress**: Most widely used, supports path-based routing and TLS.
  - **Traefik**: Lightweight, dynamic configuration, suitable for microservices.

# Path-Based Routing with NGINX Ingress example:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: api-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - host: myapp.example.com
    http:
      paths:
      - path: /api
        pathType: Prefix
      backend:
        service:
          name: api-service
          port:
            number: 80
      - path: /dashboard
        pathType: Prefix
      backend:
        service:
          name: dashboard-service
          port:
            number: 80
```



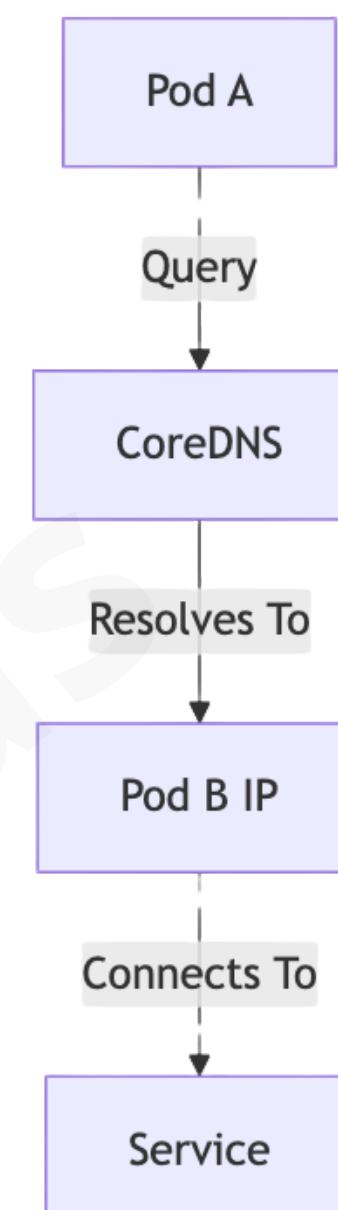
# Networking Deep Dive

## DNS & Service Discovery

Kubernetes uses CoreDNS for internal service discovery and DNS resolution.

### How Service Discovery Works

- Pods talk to each other via **service-name.namespace.svc.cluster.local**.
- **CoreDNS** automatically maps Service names to Pod IPs.



### Troubleshooting Connectivity

Check if CoreDNS is running:

```
kubectl get pods -n kube-system -l k8s-app=kube-dns
```

Test DNS resolution inside a Pod:

```
kubectl run busybox --image=busybox -it --restart=Never -- nslookup api-service.default.svc.cluster.local
```

Common Issues:

DNS not resolving? Restart CoreDNS:

```
kubectl rollout restart deployment coredns -n kube-system
```

Pod can't reach a service? Check if the service exists:

```
kubectl get svc
```



@Sandip Das

# Storage Solutions

Kubernetes storage ensures data persistence for stateful applications.

This deep dive covers:

- PersistentVolumes (PVs): Static vs. Dynamic Provisioning
- StorageClasses (EBS, GCP Persistent Disk)
- Stateful Apps & VolumeClaimTemplates
- Gotcha: Avoiding Data Loss During Pod Rescheduling

## PersistentVolumes (PV) & PersistentVolumeClaims (PVC)

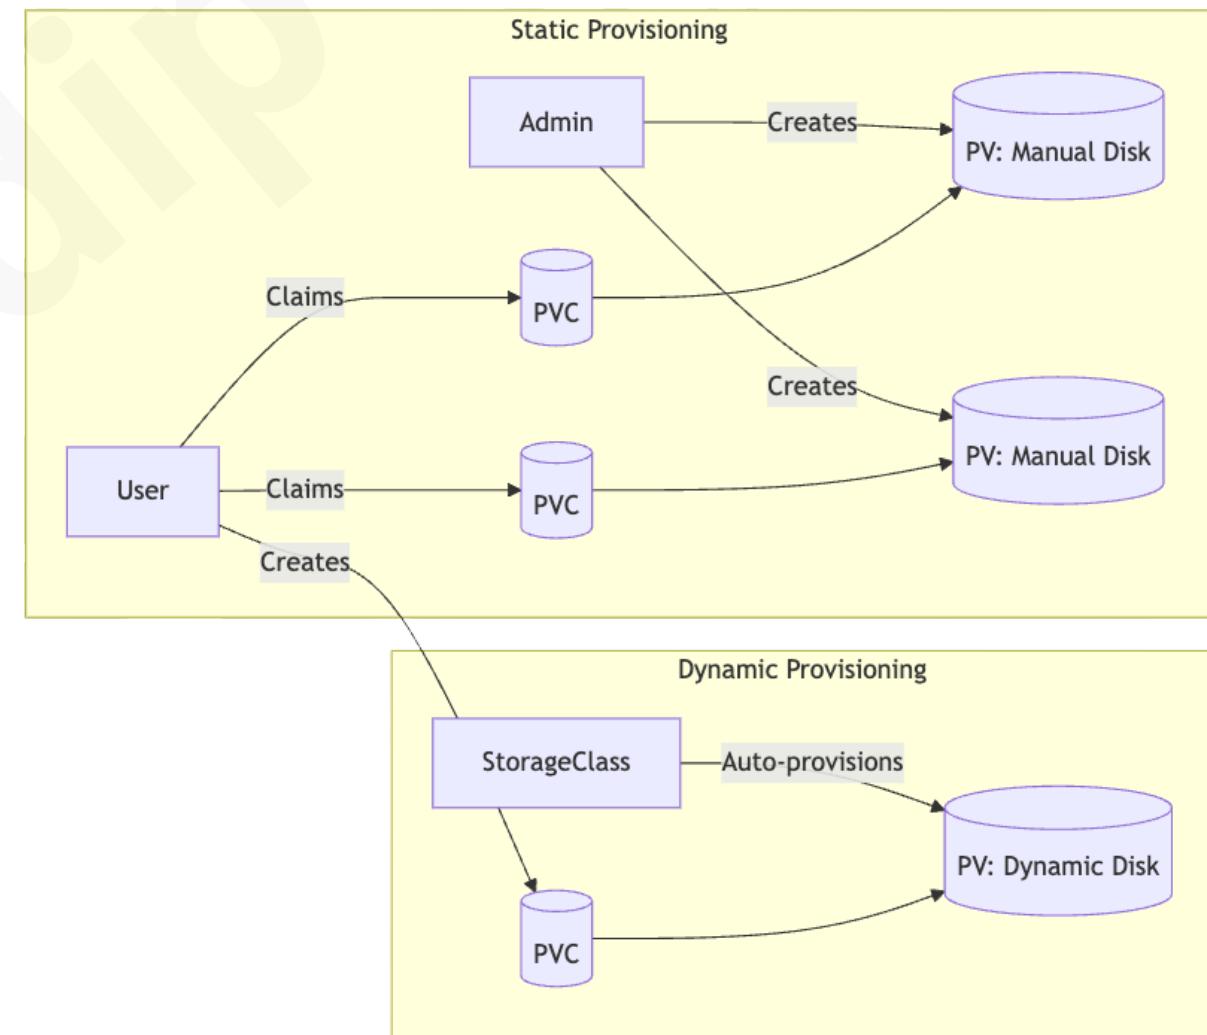
In Kubernetes,

PersistentVolumes (**PVs**) are **storage resources** independent of Pods, while PersistentVolumeClaims (**PVCs**) allow Pods to **request storage**.

### Static vs. Dynamic Provisioning

**Static:** Admin manually creates PVs, and users claim them via PersistentVolumeClaims (PVCs).

**Dynamic:** Automatically provisions PVs on-demand using StorageClass.



### Static Provisioning

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: static-pv
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/mnt/data"
```

### Dynamic Provisioning

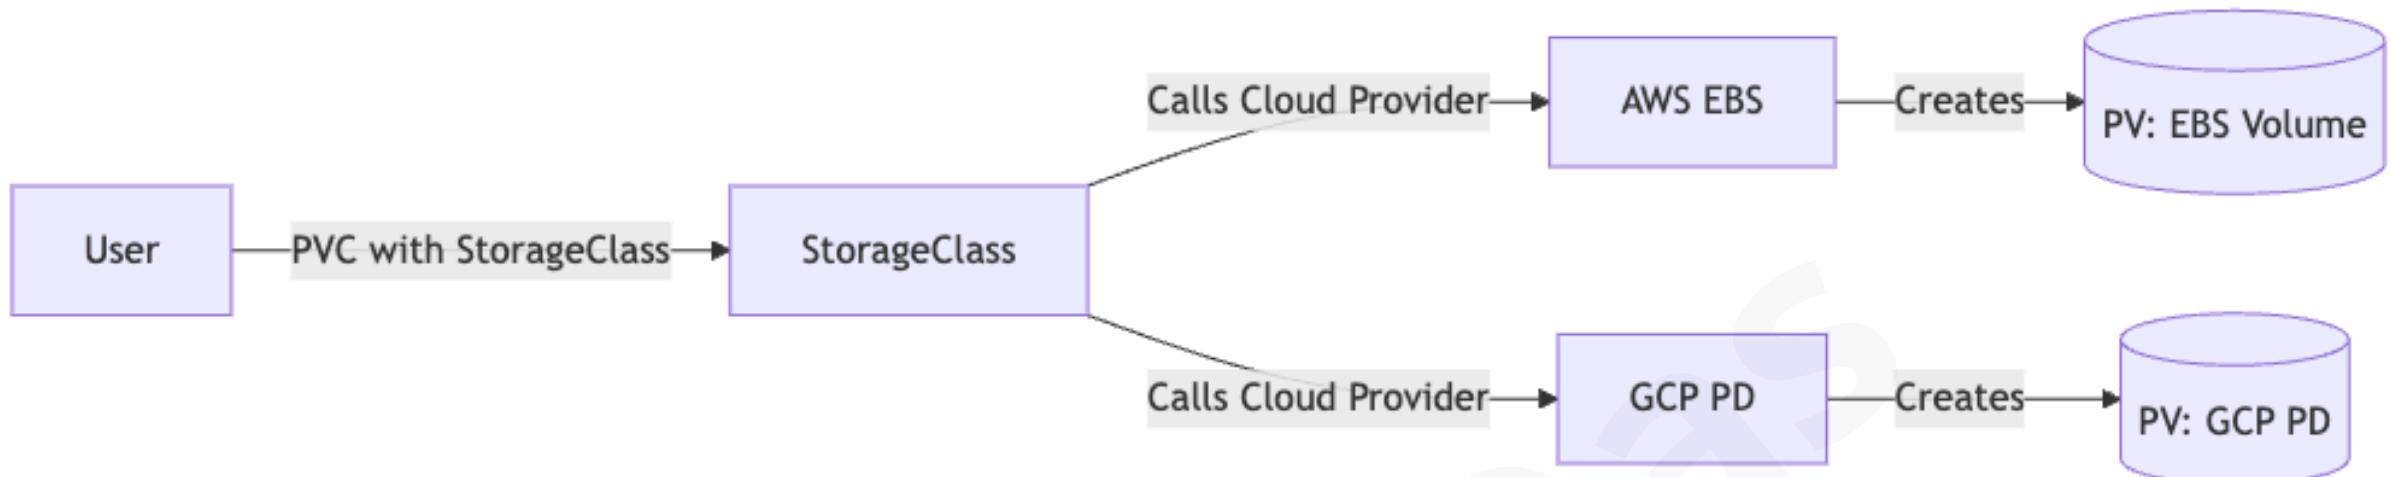
```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: dynamic-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
  storageClassName: gp2
```



@Sandip Das

# Storage Solutions

## StorageClasses: Cloud Storage Solutions



**StorageClasses enable dynamic provisioning for cloud storage:**

### Key Features:

- **Provisioner:** Defines the backend storage provider.
- **Reclaim Policy:** (Retain, Delete, Recycle).
- **Access Modes:** (ReadWriteOnce, ReadWriteMany).

### Cloud ProviderStorageClass Example

AWS (EBS) : provisioner: **ebs.csi.aws.com**

GCP (PD) : provisioner: **pd.csi.storage.gke.io**

Azure (Disk): provisioner: **disk.csi.azure.com**

#### AWS EBS StorageClass

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: gp2
provisioner: kubernetes.io/aws-ebs
parameters:
  type: gp2
  fsType: ext4
```

#### GCP Persistent Disk StorageClass

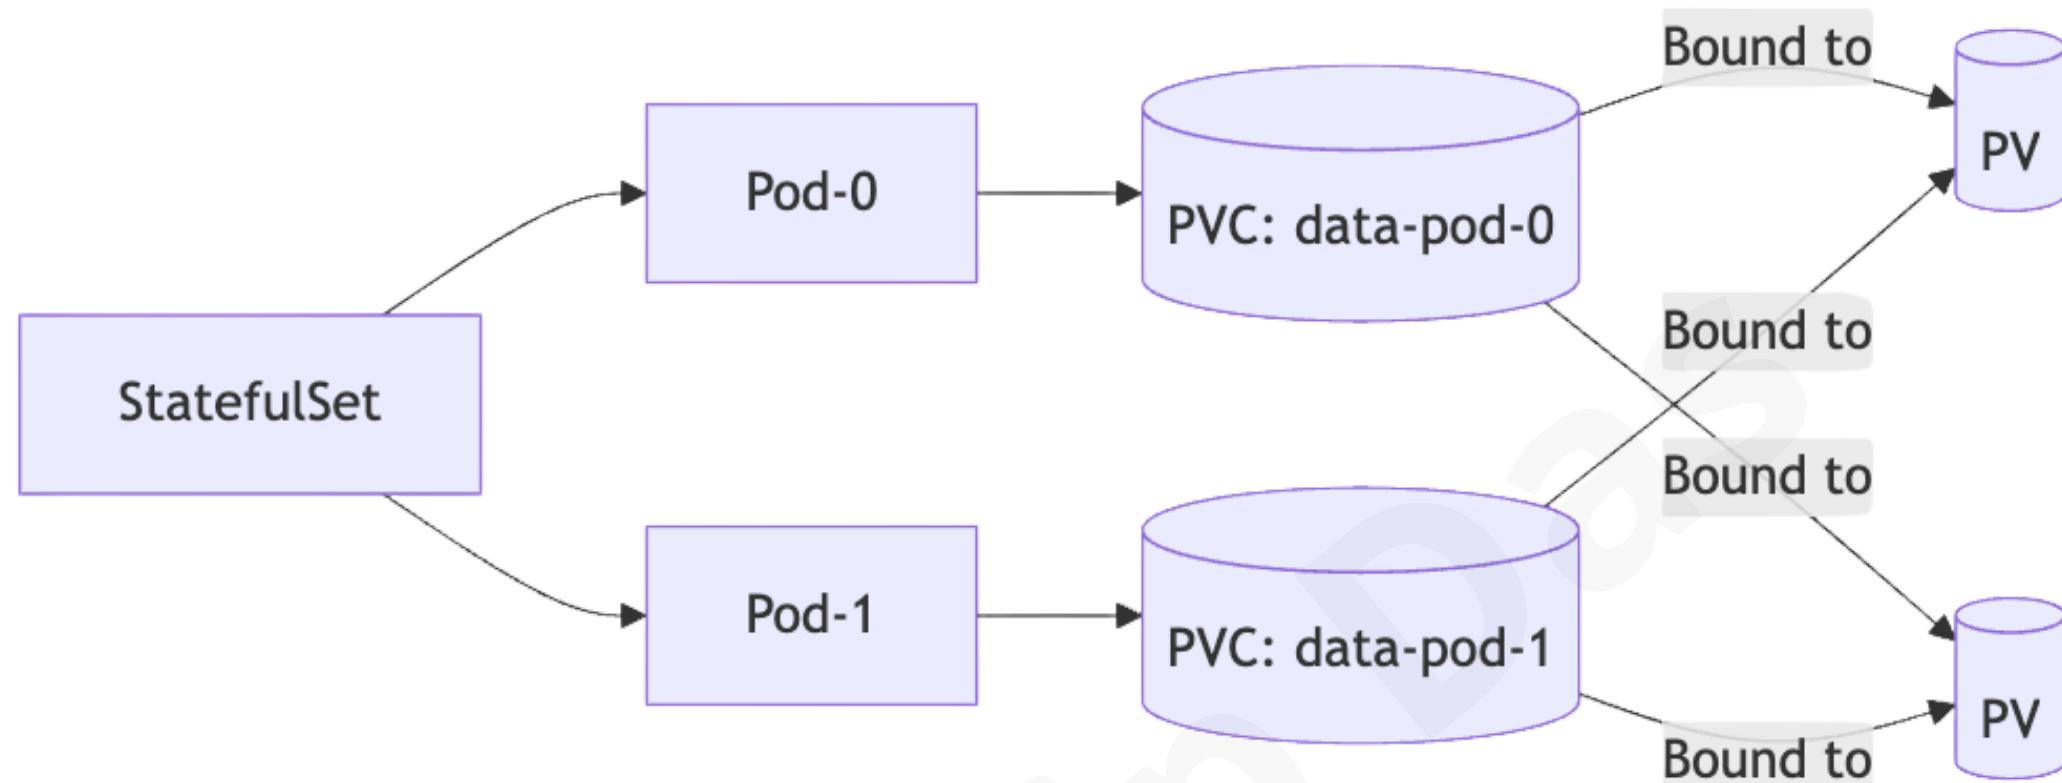
```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: standard
provisioner: kubernetes.io/gce-pd
parameters:
  type: pd-standard
  fstype: ext4
```



@Sandip Das

# Storage Solutions

## Stateful Apps & VolumeClaimTemplates in StatefulSets



**For stateful applications, Kubernetes ensures persistent storage across pod restarts and automatically creates unique PVCs for each pod in a StatefulSet (e.g., databases).**

**Example (MySQL StatefulSet):**

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: mysql
spec:
  serviceName: mysql
  replicas: 3
  selector:
    matchLabels:
      app: mysql
  template:
    metadata:
      labels:
        app: mysql
    spec:
      containers:
        - name: mysql
          image: mysql:8.0
          volumeMounts:
            - name: mysql-data
              mountPath: /var/lib/mysql
  volumeClaimTemplates:
    - metadata:
        name: mysql-data
      spec:
        accessModes: [ "ReadWriteOnce" ]
        storageClassName: "ebs-sc" #
        dynamic provisioning
        resources:
          requests:
            storage: 20Gi
```

## Why VolumeClaimTemplates?

- Each Pod in a StatefulSet gets a unique PVC.
- Prevents data corruption in distributed databases.



@Sandip Das

# Storage Solutions

## Gotcha: Avoiding Data Loss During Pod Rescheduling

### Common Issues

- Pods get rescheduled to another node, losing attached storage.
- PVCs fail to reattach, causing application failures.

### Key Strategies:

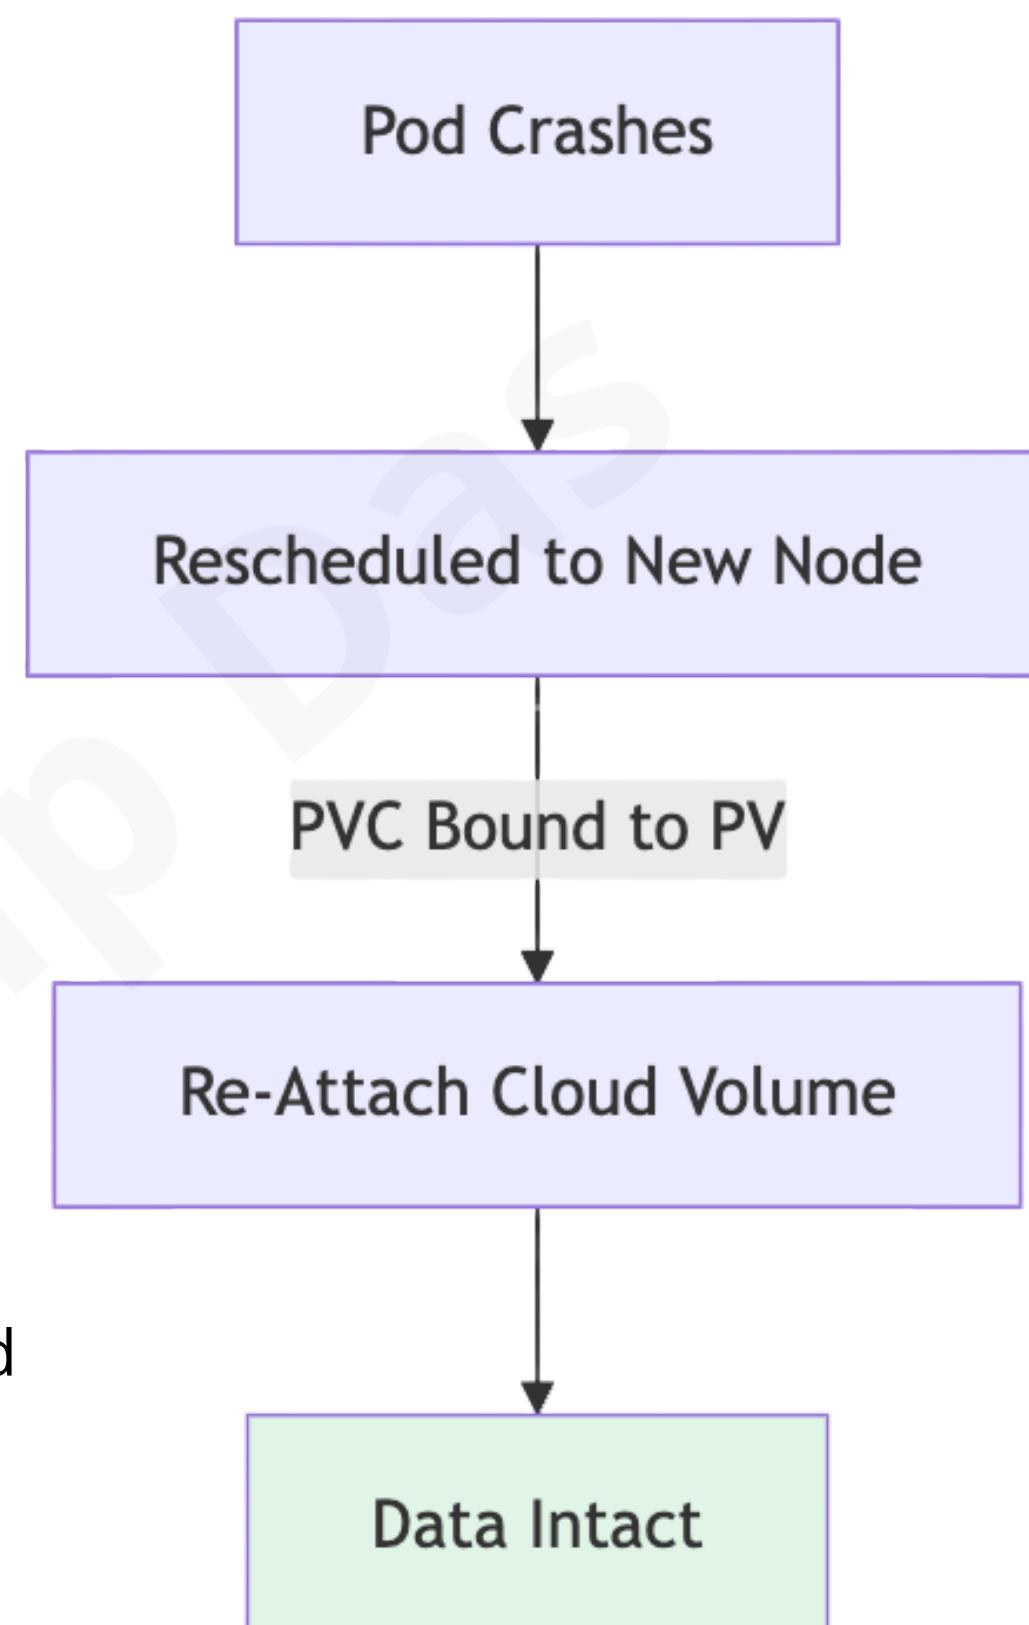
1. **Reclaim Policy:** Set PV's persistentVolumeReclaimPolicy to Retain (not Delete).
2. **Backups:** Use tools like Velero to backup PVs.
3. **Volume Binding:** Use WaitForFirstConsumer to delay PV binding until pod scheduling.
4. **ReadWriteMany (RWO):** For shared volumes (e.g., NFS), but avoid for databases.

### Troubleshooting Tips

- **PVC Pending:** Check StorageClass, resource quotas, or cloud provider limits.
- **Access Modes:** Ensure pods don't request ReadWriteOnce volumes in multi-pod deployments.
- **Backup/Restore:** Use Velero to clone PVs across clusters.

### Best Practices:

- **For AWS/GCP:** Use EBS/Persistent Disks with StorageClasses.
- **For NFS/CephFS:** Ensure multi-node attach support.



### Example (Retain Policy):

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-retain
spec:
  persistentVolumeReclaimPolicy: Retain # Critical!
  capacity:
    storage: 10Gi
  # ... rest of PV config
```



@Sandip Das

# Troubleshooting and Error Handling

## Common Errors & Fixes

### ImagePullBackOff

**Cause:** Kubernetes can't pull the image due to an incorrect image name, missing authentication, or registry issues.

#### 🔍 Troubleshooting Steps

`kubectl get pods`

`kubectl describe pod <pod-name>`

#### ✓ Fixes

⚠ Wrong Image Name? Check the exact image name:\

`containers:`

`- name: my-app`

`image: nginx:latest # Ensure the correct version/tag`

⚠ Private Registry? Create a Secret and reference it in the Pod:

`kubectl create secret docker-registry regcred \
--docker-server=<REGISTRY> --docker-username=<USER> \
--docker-password=<PASSWORD>`

In the kubernetes template

`imagePullSecrets:`

`- name: regcred`

### CrashLoopBackOff

**Cause:** The container is crashing repeatedly due to liveness probe failures, Application crashes, missing dependencies, or resource constraints (**OOMKilled**).

#### 🔍 Troubleshooting Steps:

`kubectl describe pod <pod-name>`

`kubectl logs <pod-name> -c <container-name>`

#### ✓ Fixes:

Check Application Logs (`kubectl logs`) for errors.

Increase CPU/Memory Limits if running out of resources.

`livenessProbe:`

Fix Liveness Probe Issues:

`httpGet:`

`path: /health`

`port: 8080`

`initialDelaySeconds: 5`

`periodSeconds: 10`

`resources:`

`requests:`

`memory: "256Mi"`

`cpu: "500m"`

`limits:`

`memory: "512Mi"`

`cpu: "1"`



# Troubleshooting and Error Handling

## NetworkPluginNotReady

**Cause:** Kubernetes can't communicate between Pods due to CNI misconfiguration and / or Misconfigured kubelet or cni.conf.

### 🔍 Troubleshooting Steps:

```
kubectl get nodes -o wide  
kubectl describe node <node-name>  
kubectl get pods -n kube-system
```

### ✓ Fixes

#### Check if the CNI Plugin is Installed:

```
ls /etc/cni/net.d/
```

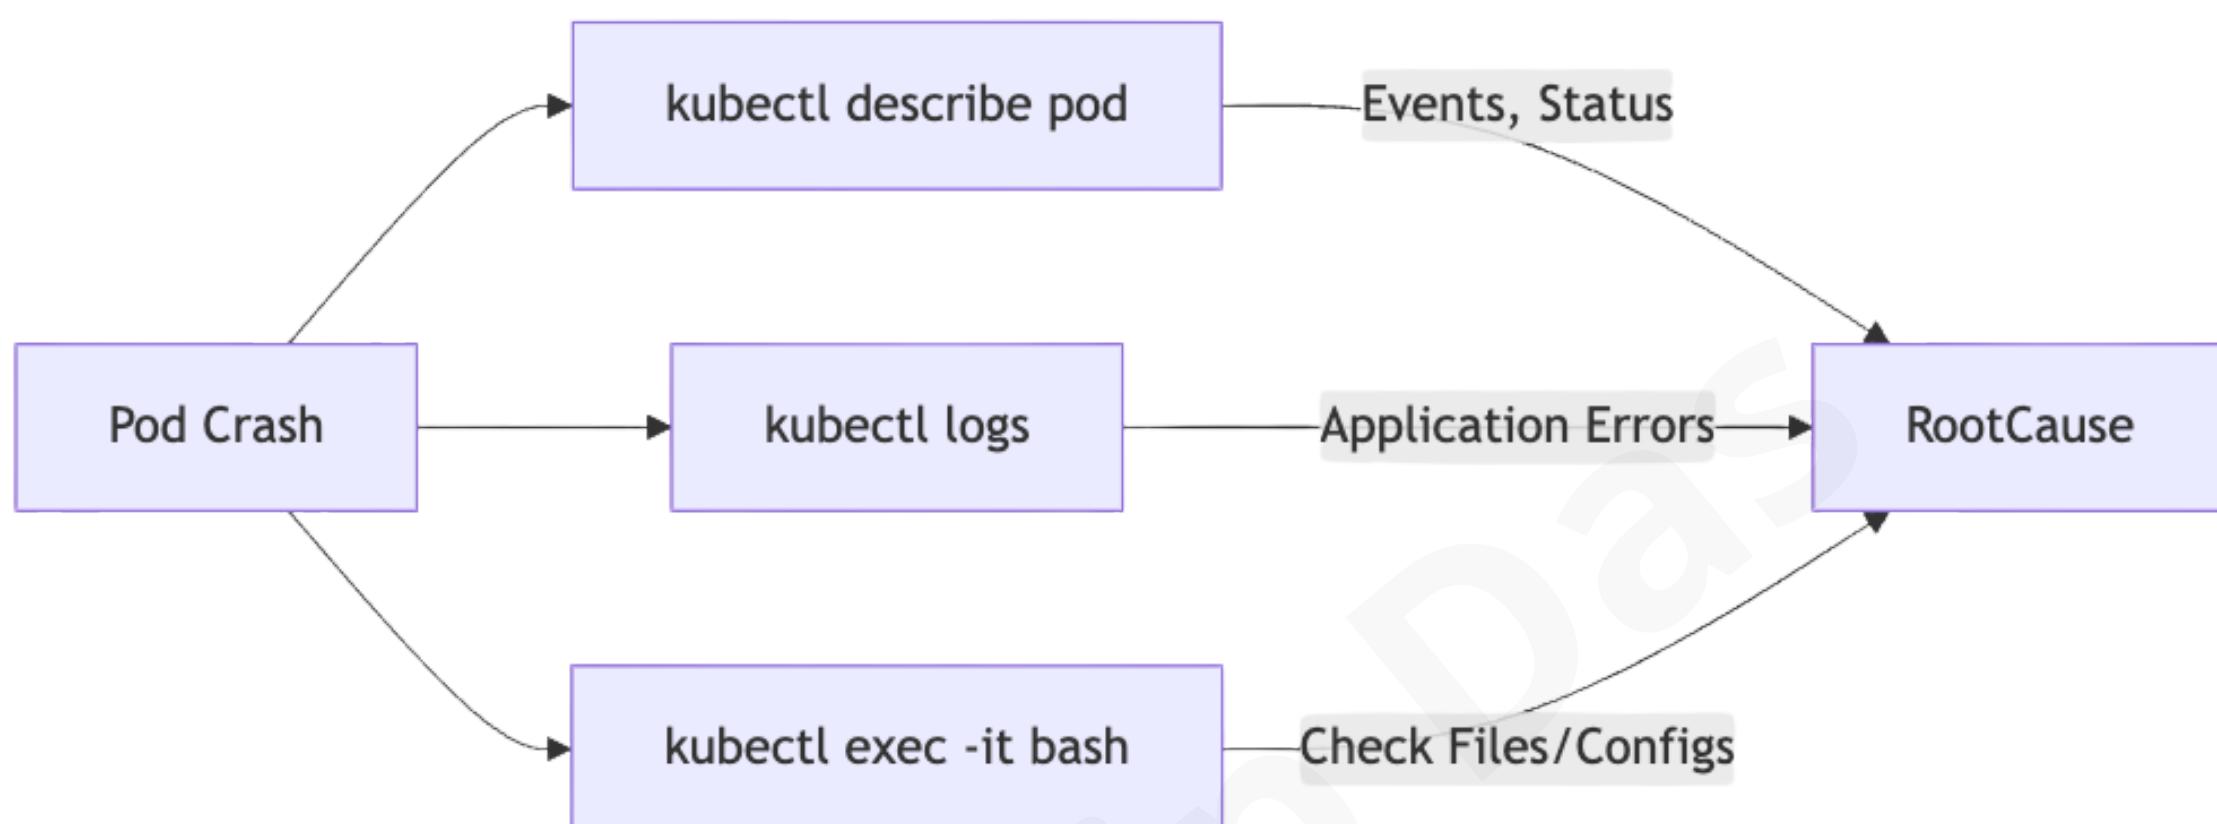
#### Reinstall Flannel/Calico CNI Plugin:

```
kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
```



# Troubleshooting and Error Handling

## Debugging Tools



### 🔍 1 kubectl describe (Detailed Pod/Node Info):

`kubectl describe pod <pod-name>`  
`kubectl describe node <node-name>`

✓ Use it for:

- Event logs, failed mounts, and scheduling failures.

### 🔍 2 kubectl logs (Container Logs):

`kubectl logs <pod-name> -c <container-name>`

✓ Use it for:

- Checking application errors.
- Seeing crash reasons (e.g., missing dependencies).

### 🔍 3 kubectl exec (Debugging Inside a Pod):

`kubectl exec -it <pod-name> -- /bin/sh`

✓ Use it for:

- Checking files inside a container (ls, cat /var/log/app.log).
- Testing database connections (mysql -h db-service -u root -p).



@Sandip Das

# Troubleshooting and Error Handling

## Real-World Scenarios & Fixes

### Scenario 1: Service Endpoints Missing (Label Mismatch)

**Issue:** A Pod is running, but the Service can't route traffic to it.

#### 🔍 Check if Endpoints Exist:

kubectl get endpoints <service-name>

#### ✓ Fix

- Ensure Pod Labels Match the Service Selector:

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: my-app # Must match Pod labels
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
spec:
  template:
    metadata:
      labels:
        app: my-app # Ensure it matches Service
selector
```

### Scenario 2: PersistentVolume Stuck in "Pending" (Storage Class Issues)

**Issue:** A Pod requests storage, but the PersistentVolume is stuck in Pending.

#### 🔍 Check PVC Status:

kubectl get pvc

kubectl describe pvc <pvc-name>

#### ✓ Fix

Ensure a StorageClass is Defined:

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: fast-storage
provisioner: kubernetes.io/aws-ebs
```

Ensure the PV & PVC Match:

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  storageClassName: fast-storage
```



# Optimization & Cost Management

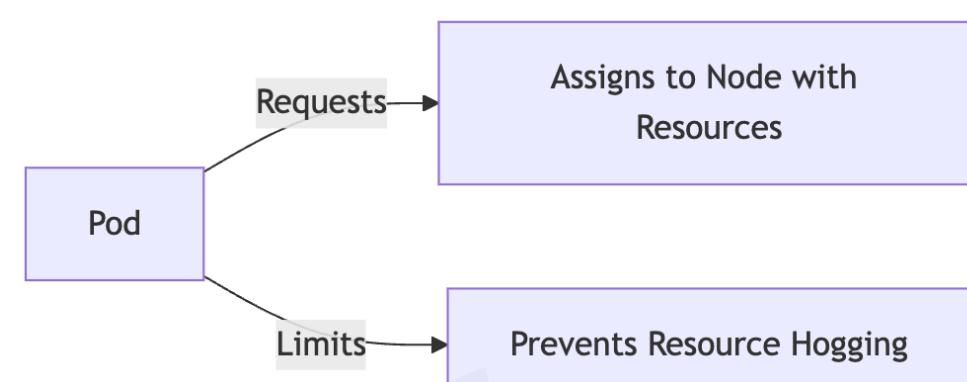
## Resource Requests & Limits:

**Purpose:** Avoid CPU throttling and OOM (Out-Of-Memory) kills by reserving and capping resources.

What Are Resource Requests & Limits?

- **Requests:** Minimum guaranteed CPU/memory for a Pod.
- **Limits:** Maximum CPU/memory a Pod can use.

```
apiVersion: v1
kind: Pod
metadata:
  name: optimized-app
spec:
  containers:
  - name: app
    image: nginx
    resources:
      requests: # Guaranteed minimum
        cpu: "100m"
        memory: "256Mi"
      limits: # Maximum allowed
        cpu: "500m"
        memory: "512Mi"
```



Issue	Cause	Fix
CPU Throttling	CPU requests are too low	Increase CPU requests
OOM Kills	Memory limit exceeded	Set proper memory limits
Node Overcommit	Requests exceed node capacity	Adjust pod placement

## Key Tips:

1. Always set requests to match your app's baseline needs.
2. Use limits to prevent runaway processes from starving other pods.
3. Monitor with kubectl top pods or tools like Prometheus.



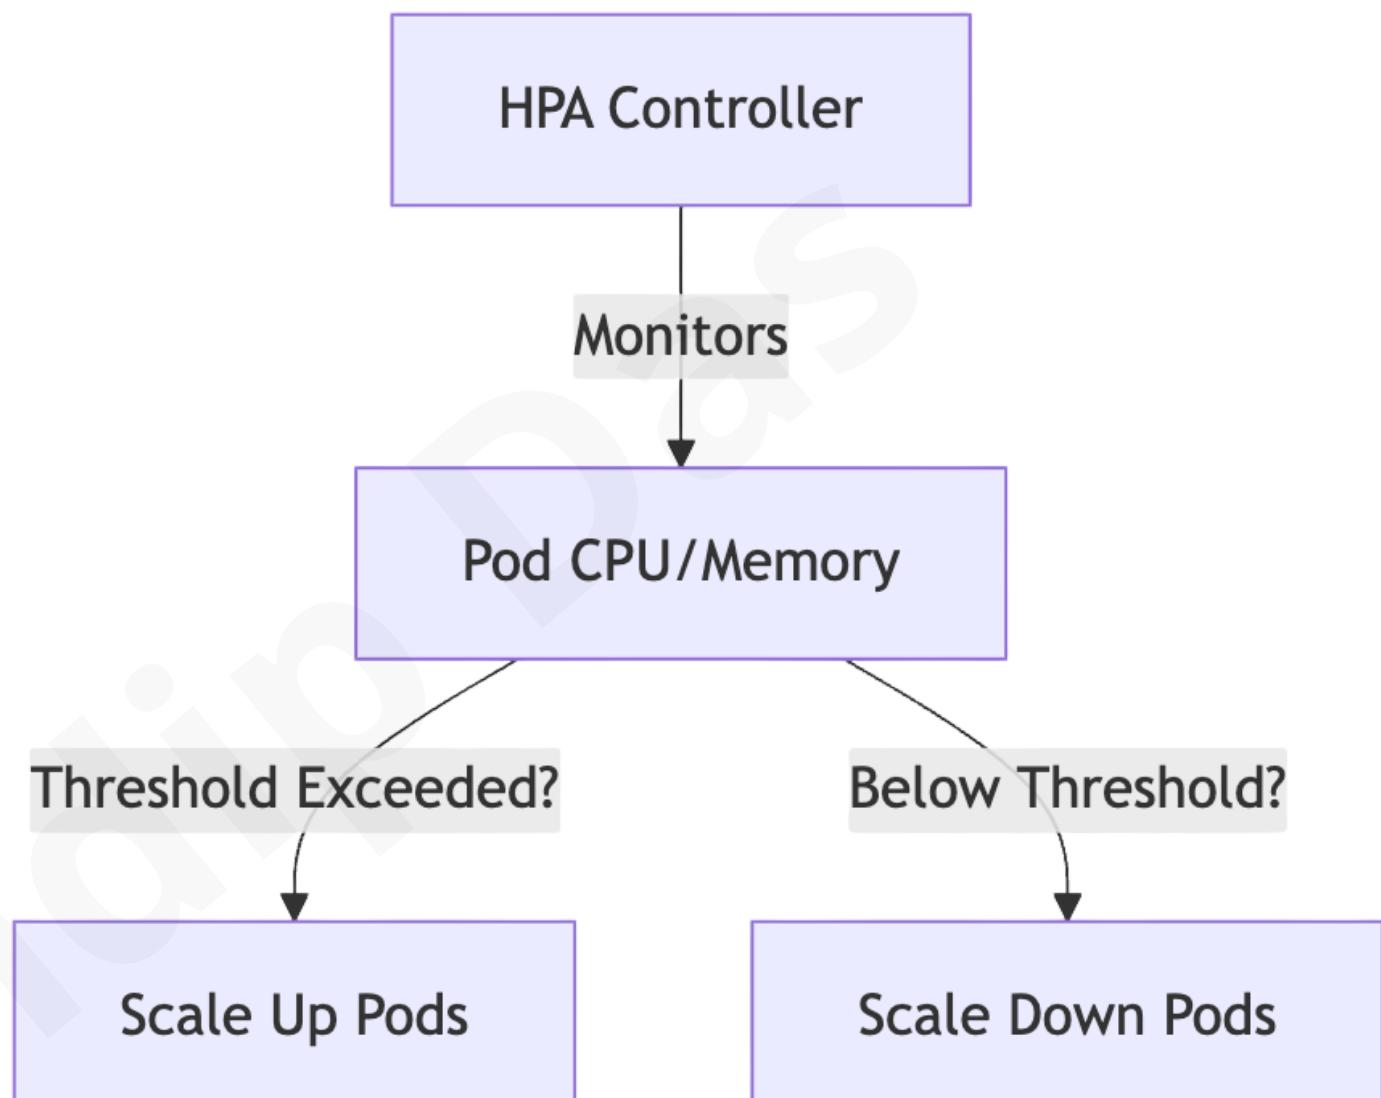
# Optimization & Cost Management

## Horizontal Pod Autoscaler (HPA):

HPA automatically scales Pods based on CPU, memory, or custom metrics.

### Example: Scale Pods Based on CPU

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: my-app-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: my-app
  minReplicas: 2
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 70
```



### Key Tips:

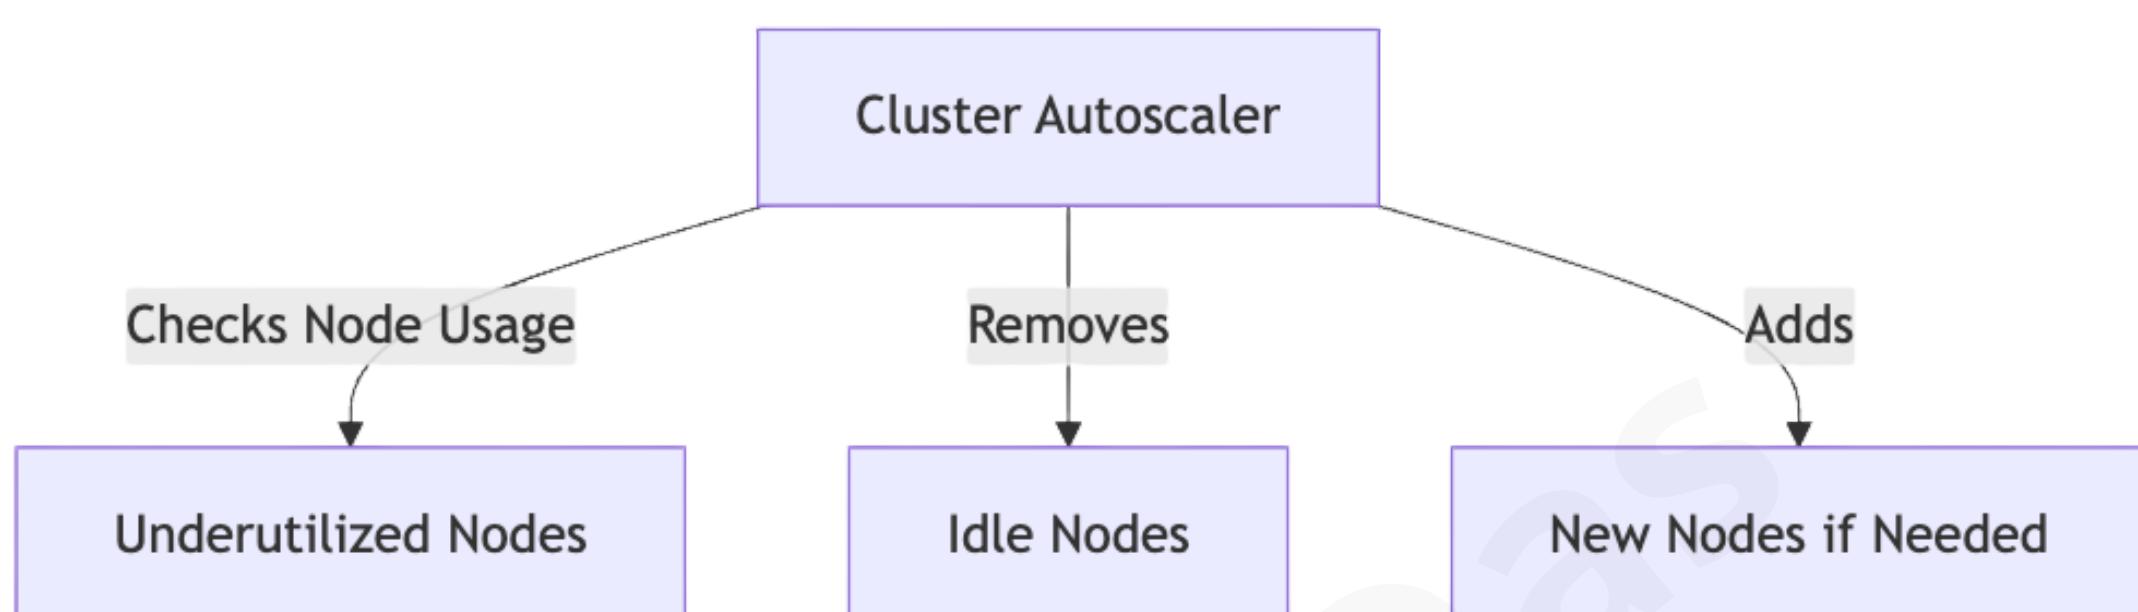
1. Use custom metrics (e.g., requests per second) for precise scaling.
2. Combine with Cluster Autoscaler to scale nodes when needed.



@Sandip Das

# Optimization & Cost Management

## Cluster Autoscaler (CA)



**Purpose:** Automatically **resize (add/remove nodes or VMs)** node pools based on pod scheduling demands.

**Enable Cluster Autoscaler (EKS/GKE/AKS):**  
`kubectl apply -f cluster-autoscaler.yaml`

### cluster-autoscaler.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: cluster-autoscaler
  namespace: kube-system
data:
  scale-down-enabled: "true"
  scale-down-unneeded-time: "10m"
  scale-down-utilization-threshold: "0.5"
```

### Key Tips:

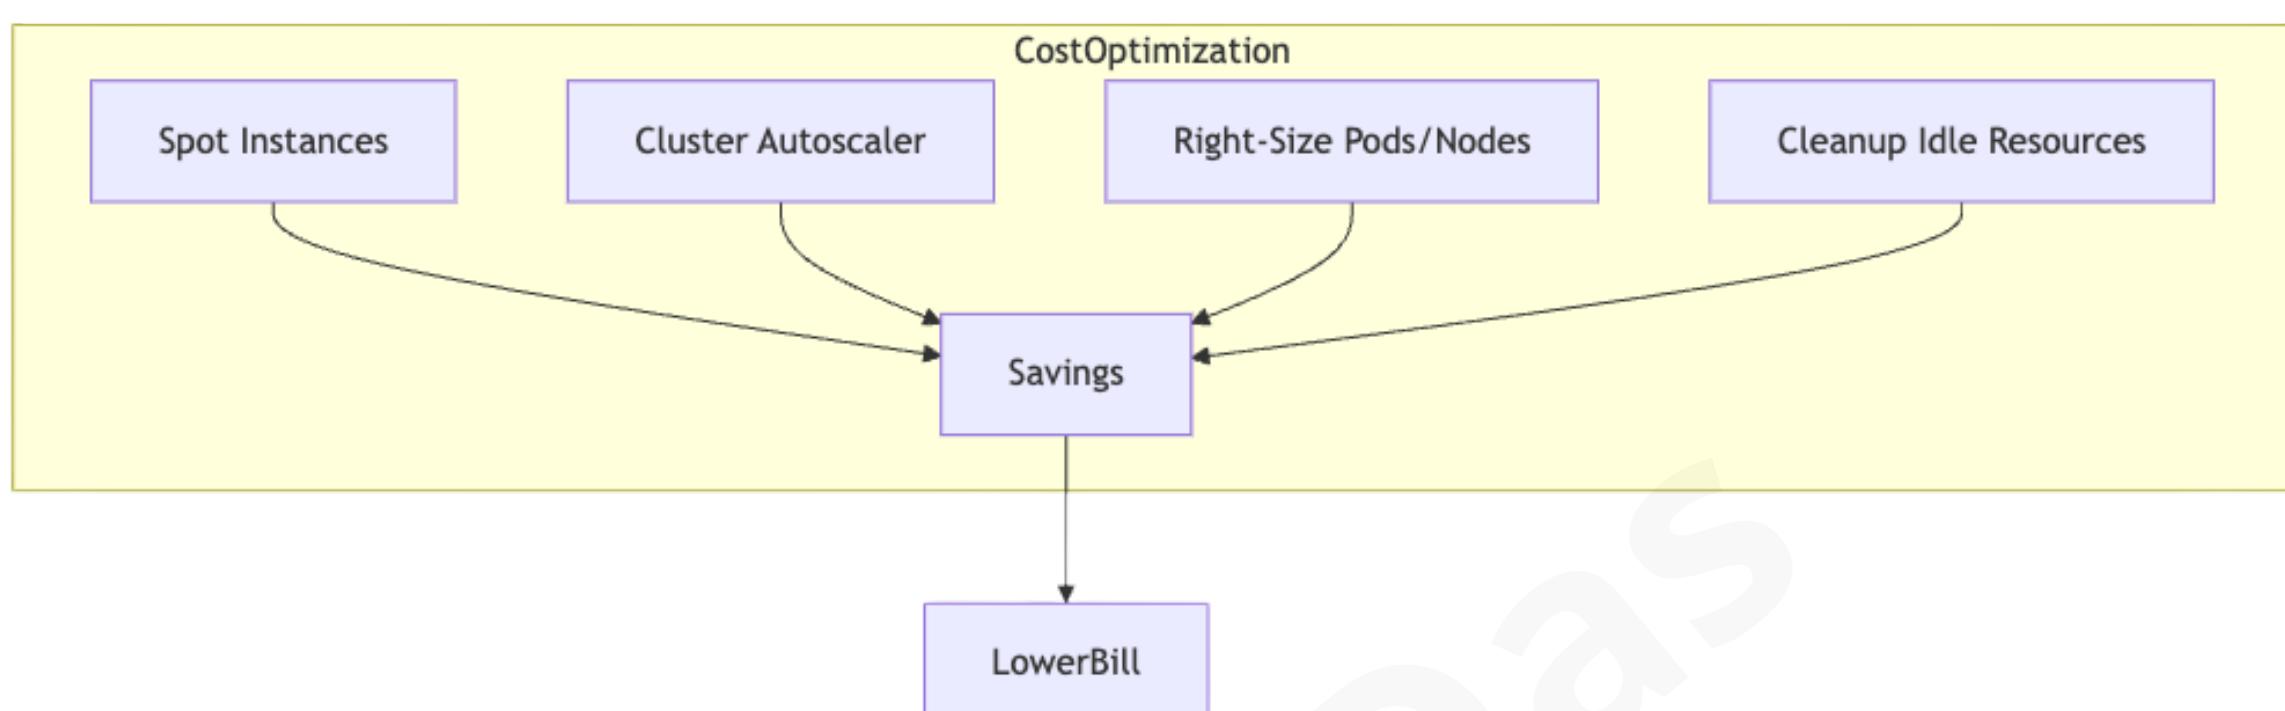
1. **Configure node groups with mixed instance types (e.g., spot + on-demand).**
2. **Use pod priority to prioritize critical workloads during scaling.**



@Sandip Das

# Optimization & Cost Management

## Cost Saving Tips



### Spot Instances (AWS/GCP/Azure)

- Use for stateless, fault-tolerant workloads (e.g., batch jobs).
- Save 60–90% compared to on-demand instances.

### Avoid Overprovisioning

- Right-size node pools (avoid oversized VMs).
- Use Vertical Pod Autoscaler (VPA) to adjust requests/limits automatically.

### Optimize Pod Density

- Pack more pods per node (balance with resource limits).
- Use Karpenter (AWS) for cost-aware node provisioning.

### Cleanup Idle Resources

- Delete unused pods, services, and PVs.
- Schedule preemptible/spot nodes for non-critical workloads.

### Real-World Example: Cost-Optimized Workflow

1. **Stateless Apps:** Deploy on spot instances with HPA.
2. **Stateful Apps:** Use on-demand + regular snapshots.
3. **Batch Jobs:** Run on spot instances with retries.

### Tools for Monitoring Costs

- **Kubecost:** Track cluster spend by namespace/deployment.
- **Cloud Provider Tools:** AWS Cost Explorer, GCP Cost Management.



# Kubernetes Gotchas and Anti-Patterns

While Kubernetes is powerful, certain gotchas and anti-patterns can lead to downtime, security vulnerabilities, and cost inefficiencies. Now we will cover:

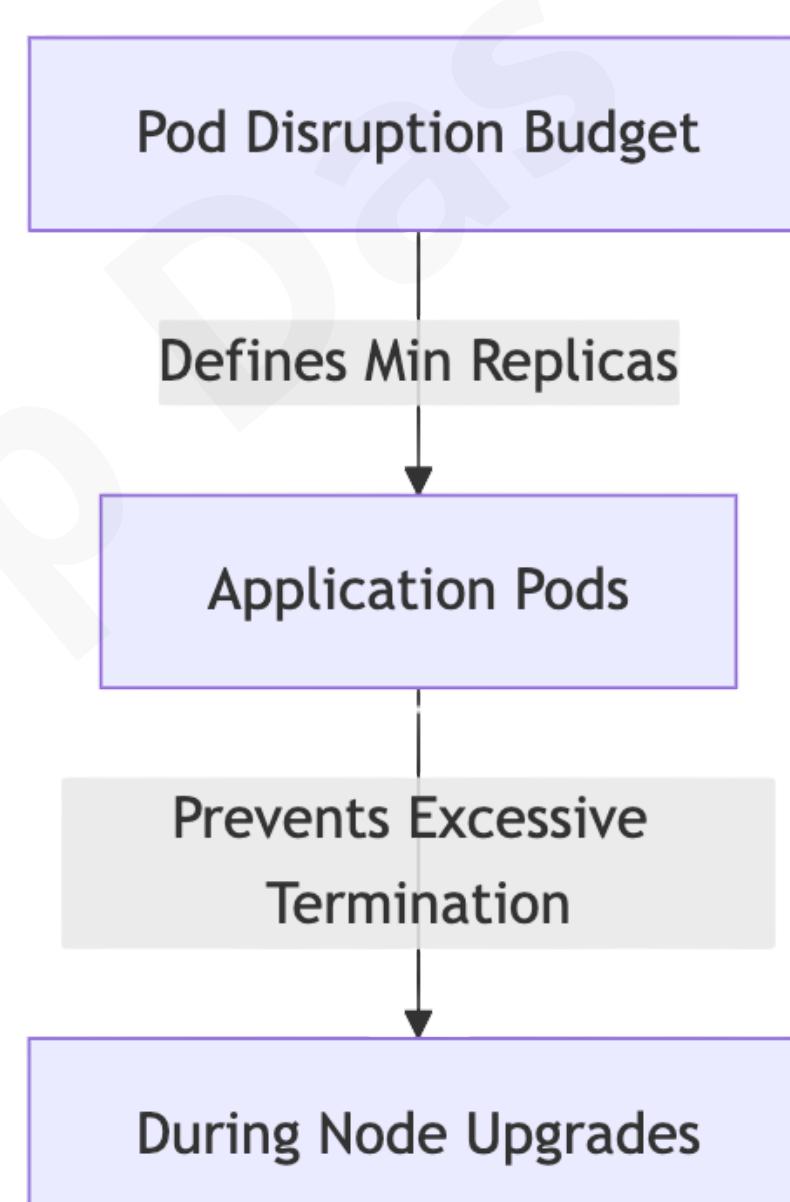
- Pod Disruption Budgets (PDBs): Avoiding unnecessary downtime during upgrades
- HostPath Volumes: Security risks and safer alternatives
- Zombie Namespaces: Orphaned resources silently consuming costs
- Version Skew: Control plane and worker node version mismatches

## Pod Disruption Budgets (PDBs)

**Problem:** Upgrades/rollouts can cause unintended downtime if PDBs are not set to protect critical pods.

### Example: PDB for High Availability

```
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: my-app-pdb
spec:
  minAvailable: 2 # At least 2 pods
  mustBeRunning
  selector:
    matchLabels:
      app: my-app
```



**Pro Tip:**  
**Always define PDBs for stateful or high-availability workloads (e.g., databases, API gateways).**



@Sandip Das

# Kubernetes Gotchas and Anti-Patterns

## HostPath Volumes: Security Risks

hostPath volumes mount directories from the host node directly into a Pod.

**Problem:** This bypasses Kubernetes' abstraction and can introduce security risks as it mounts expose host filesystems, risking privilege escalation or data corruption.

### Anti-Pattern: (BAD)

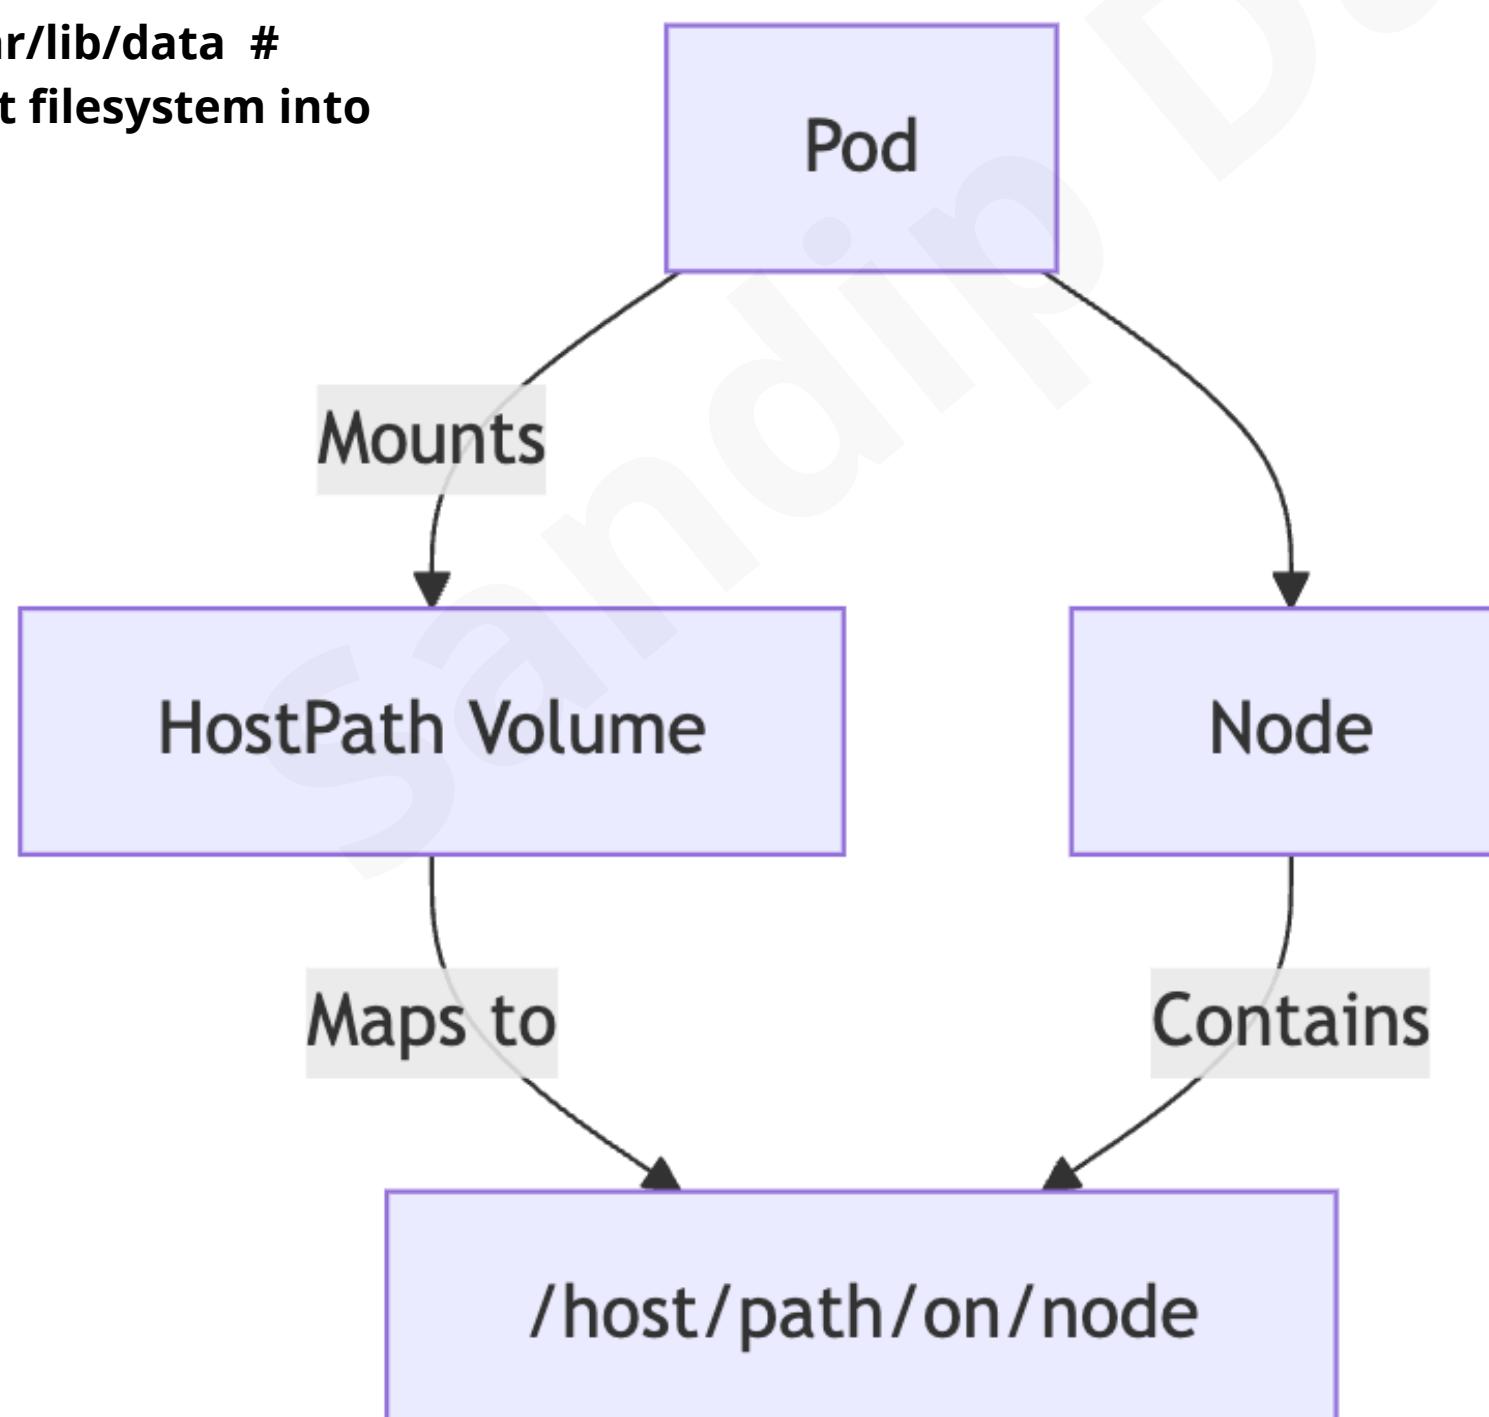
```
apiVersion: v1
kind: Pod
metadata:
  name: hostpath-pod
spec:
  containers:
    - name: my-container
      image: nginx
      volumeMounts:
        - mountPath: /data
          name: host-data
  volumes:
    - name: host-data
      hostPath:
        path: /var/lib/data #
```

**Mounts host filesystem into container!**

### Solution: (Good)

- Use PersistentVolumes (PV) or CSI drivers for cloud storage.
- If unavoidable, restrict hostPath to read-only and specific directories:

```
volumes:
  - name: logs
    hostPath:
      path: /var/logs
    type: DirectoryOrCreate
    readOnly: true
```



# Kubernetes Gotchas and Anti-Patterns

## Zombie Namespaces

A zombie namespace is a stuck or orphaned namespace that still contains resources but isn't actively used.

**Problem:** Deleting a namespace doesn't always clean up resources (e.g., cloud load balancers, PVs), leading to lingering costs.

### How to Identify:

```
# List resources in "Terminating" namespaces  
kubectl get ns --field-selector status.phase=Terminating
```

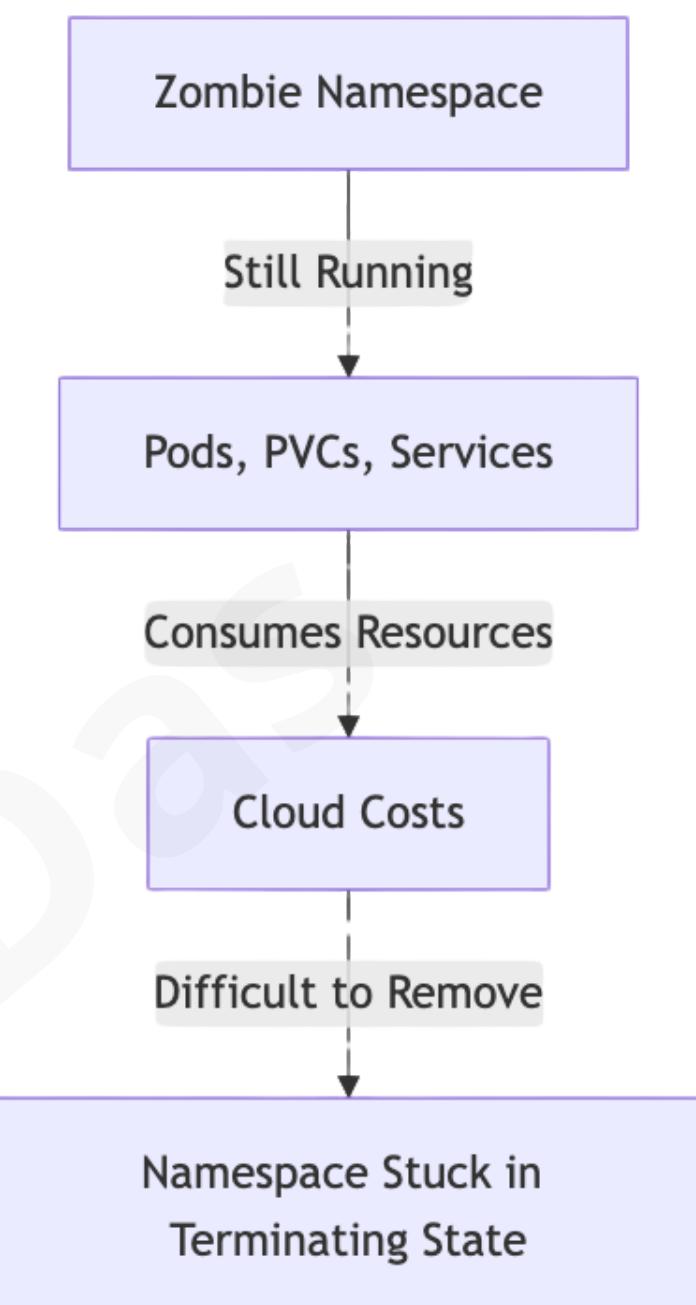
```
# Check for orphaned cloud resources (e.g., AWS EBS volumes, GCP disks).
```

### Fix:

Force-delete the namespace:

```
kubectl get namespace <name> -o json > tmp.json  
# Remove `kubernetes` from finalizers array in tmp.json  
kubectl replace --raw "/api/v1/namespaces/<name>/finalize" -f tmp.json
```

 **Pro Tip:** Use tools like Kyverno to enforce cleanup policies.

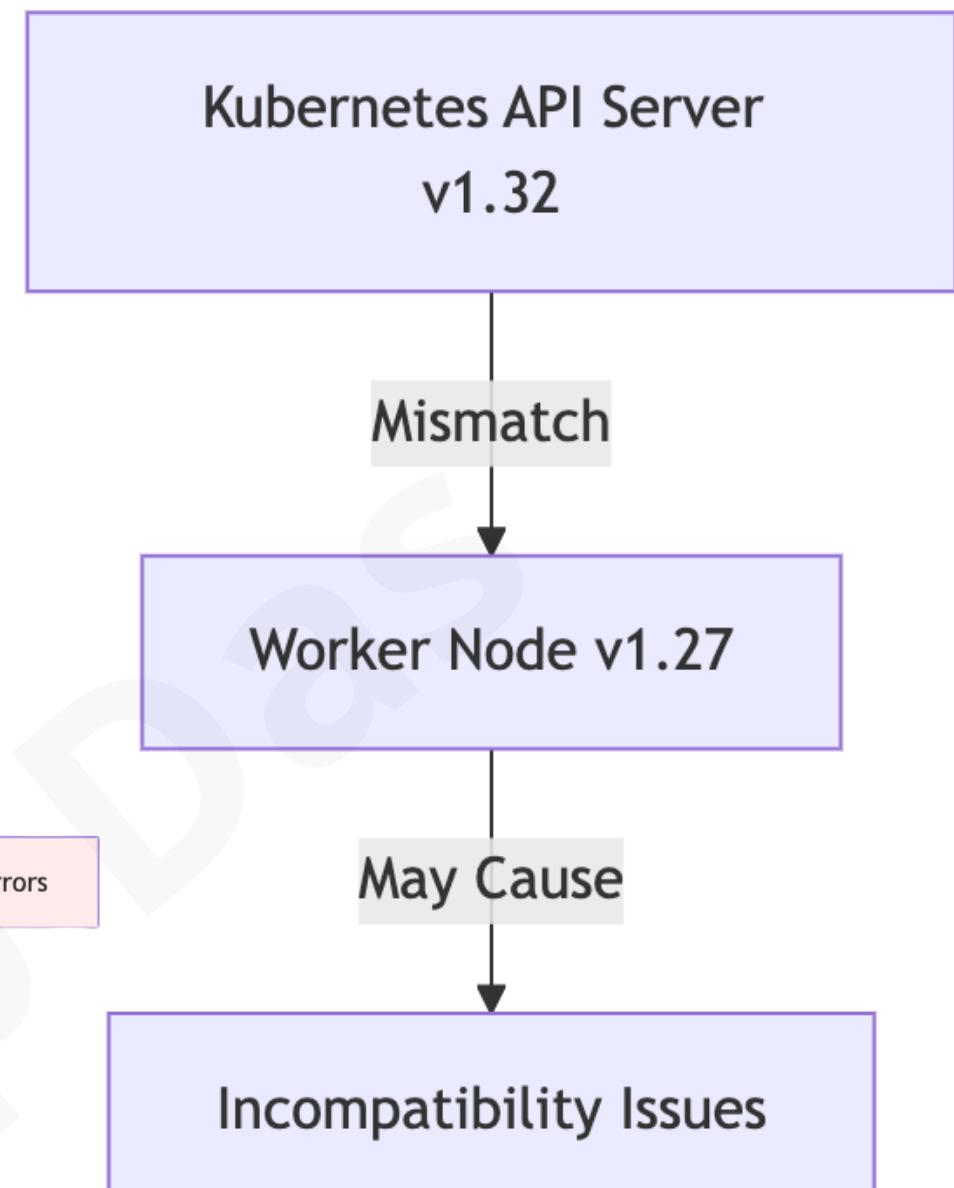
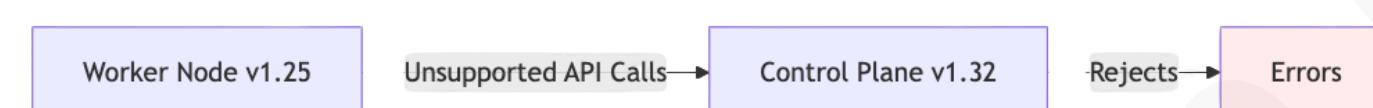


# Kubernetes Gotchas and Anti-Patterns

## Version Skew

Kubernetes releases frequent updates, and sometimes control plane and worker nodes end up running different versions, it's called **Version Skew**.

**Problem:** API deprecations or unsupported features when components are too far apart.



**Fix:**  
Check version and upgrade incrementally

```
# Check versions  
kubectl version --short  
kubelet --version
```

```
# Upgrade nodes incrementally (e.g., 1.25 → 1.26 → 1.27....1.32).
```



@Sandip Das

# Kubernetes Security Best Practices

Security in Kubernetes is critical to protect your cluster from unauthorized access, malicious workloads, and network attacks. Now we will cover:

- ✓ RBAC (Role-Based Access Control): Controlling user and service permissions
- ✓ Pod Security Admission (PSA) : Restricting privileged pods
- ✓ Network Policies: Isolating namespaces and securing traffic
- ✓ Image Security: Scanning images and avoiding vulnerabilities

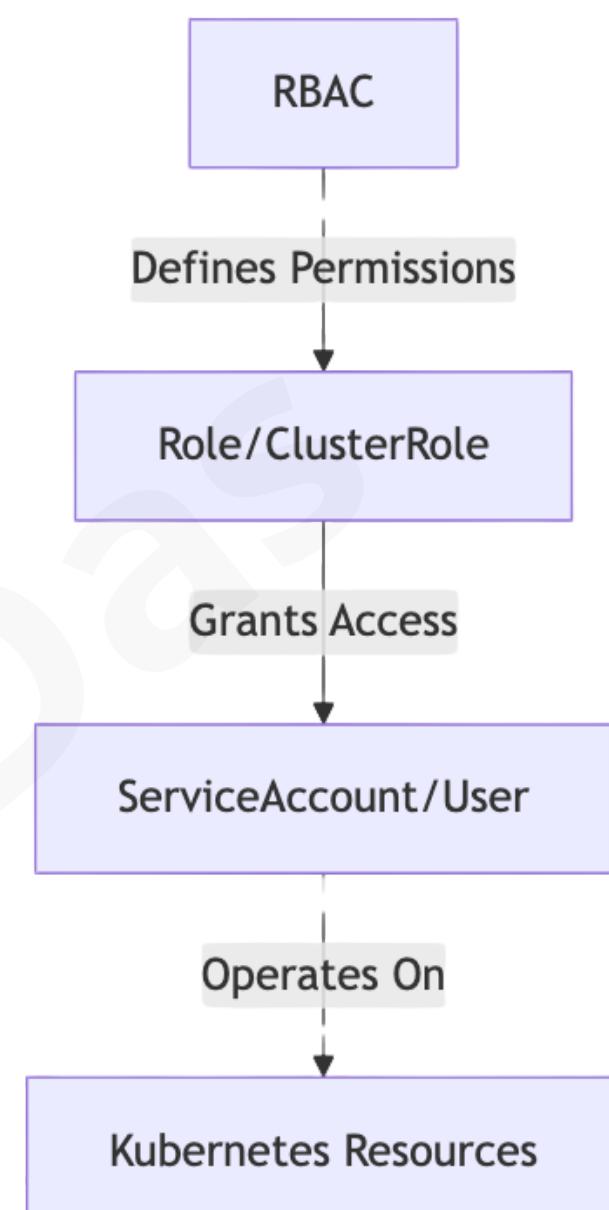
## Role-Based Access Control (RBAC)

RBAC restricts access based on Roles, ClusterRoles, and ServiceAccounts.

Example:

```
# Role: Read-only access to pods in "monitoring" namespace
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: monitoring
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "list", "watch"]
```

```
# RoleBinding: Assign to a ServiceAccount
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-pods
  namespace: monitoring
subjects:
- kind: ServiceAccount
  name: prometheus-sa
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```



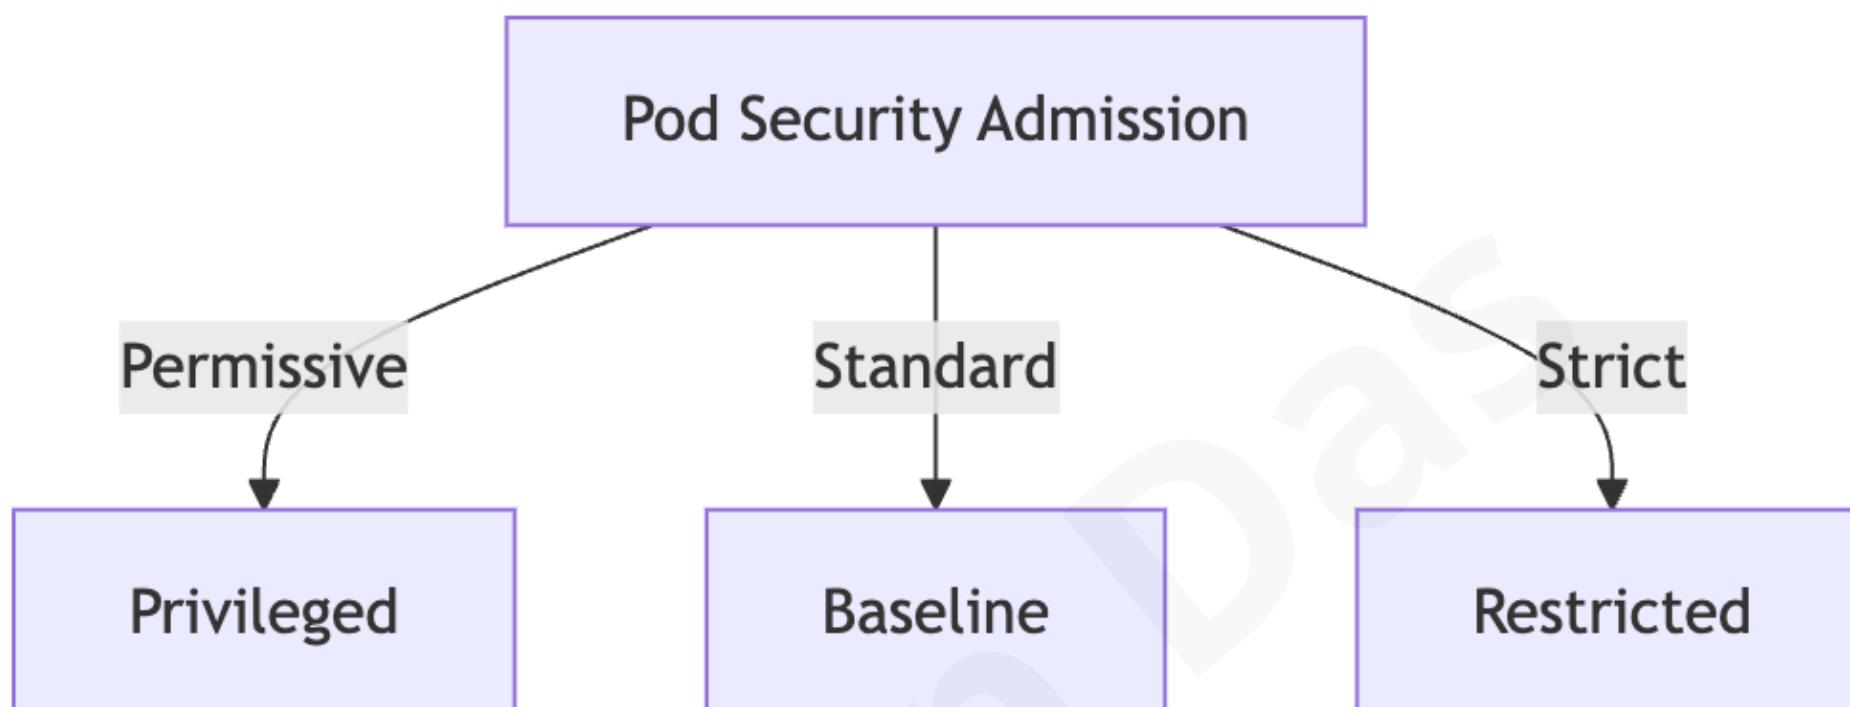
**💡 Best Practices:**  
**Avoid wildcard (\*) permissions.**  
**Use dedicated ServiceAccounts (not default).**



@Sandip Das

# Kubernetes Security Best Practices

## Pod Security Policies (PSPs)



Pod Security Admission (PSA) is the replacement for PodSecurityPolicies (PSP), which was deprecated in Kubernetes 1.21 and removed in Kubernetes 1.25. PSA enforces security policies at the namespace level using predefined security standards.

### Enforcing PSA at Namespace Level

You assign PSA levels using labels on namespaces.

```
kubectl label namespace dev pod-security.kubernetes.io/enforce=baseline  
kubectl label namespace prod pod-security.kubernetes.io/enforce=restricted
```

PSA can be applied in three modes:

**Enforce** : Rejects non-compliant Pods before creation

**Audit** : Logs warnings for non-compliant Pods

**Warn**: Displays a warning but still allows Pod creation

e.g.

```
kubectl label namespace dev pod-security.kubernetes.io/enforce=baseline
```

```
kubectl label namespace dev pod-security.kubernetes.io/audit=restricted
```

```
kubectl label namespace dev pod-security.kubernetes.io/warn=restricted
```

Example Privilege pod:

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: privileged-pod  
  namespace: prod  
spec:  
  containers:  
    - name: nginx  
      image: nginx  
      securityContext:  
        privileged: true
```

Expected Outcome:

Pod creation fails due to PSA enforcing restricted.

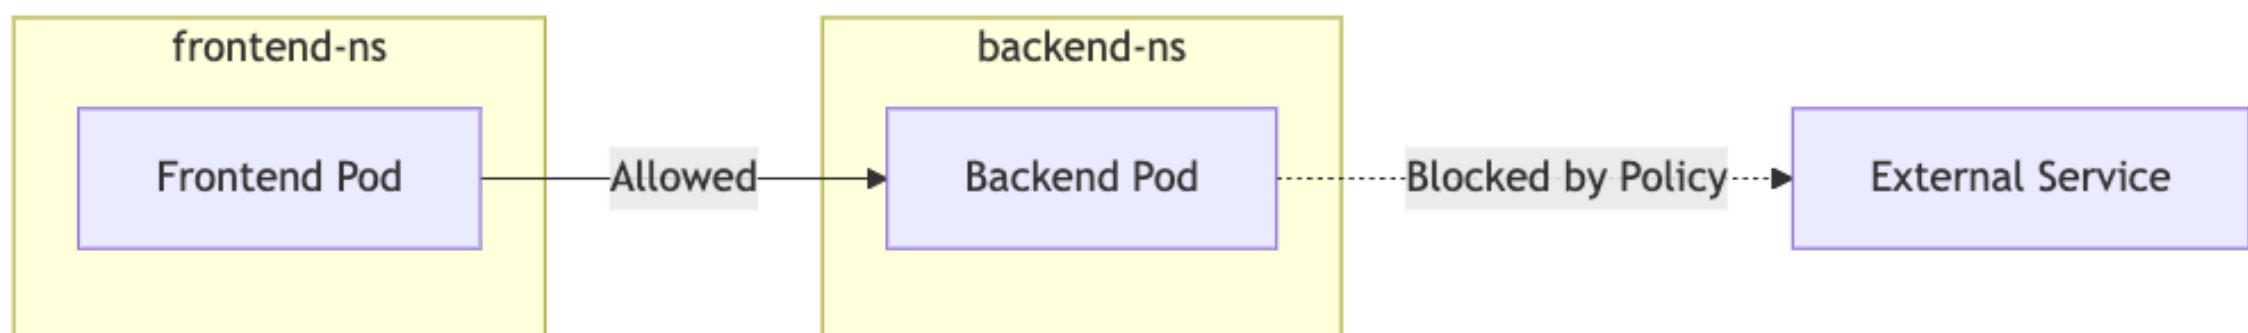
Kubernetes denies privileged containers in the prod namespace.



@Sandip Das

# Kubernetes Security Best Practices

## Network Policies



By default, Kubernetes allows all Pods to talk to each other across namespaces.

**Principle:** Deny all traffic by default, then allow specific communication.

**Example (Deny All + Allow Namespace Traffic):**

```
# Deny all ingress/egress by default
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-all
spec:
  podSelector: {}
  policyTypes:
  - Ingress
  - Egress
```

```
# Allow traffic from "frontend" namespace
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-frontend
spec:
  podSelector:
    matchLabels:
      app: backend
  ingress:
  - from:
    - namespaceSelector:
        matchLabels:
          kubernetes.io/metadata.name: frontend
```



# Kubernetes Security Best Practices

## Image Security



Using **insecure images** can expose your cluster to attacks.

**Principle:** Scan images for vulnerabilities and avoid mutable tags like :latest.

### Best Practices:

- Use immutable tags (e.g., sha256:abc123 or v1.2.3).
- Integrate tools like Trivy or Clair into CI/CD.
- Block unsigned images with Cosign/Notary.

### Example (ImagePullPolicy):

```
containers:- name: app
  image: my-registry/app:v1.2.3
  imagePullPolicy: IfNotPresent # Avoid "Always" for mutable tags
```



@Sandip Das

# Monitoring, Logging, and Observability

Observability in Kubernetes consists of metrics, logging, and tracing to gain insights into application performance and health. now we will cover:

- ✓ Prometheus/Grafana → Metrics collection & alerting
- ✓ EFK Stack (Elasticsearch, Fluentd, Kibana) → Centralized logging
- ✓ Distributed Tracing (Jaeger, OpenTelemetry) → Request flow visualization

## Prometheus/Grafana

**Prometheus** collects time-series data, and **Grafana** visualizes it.

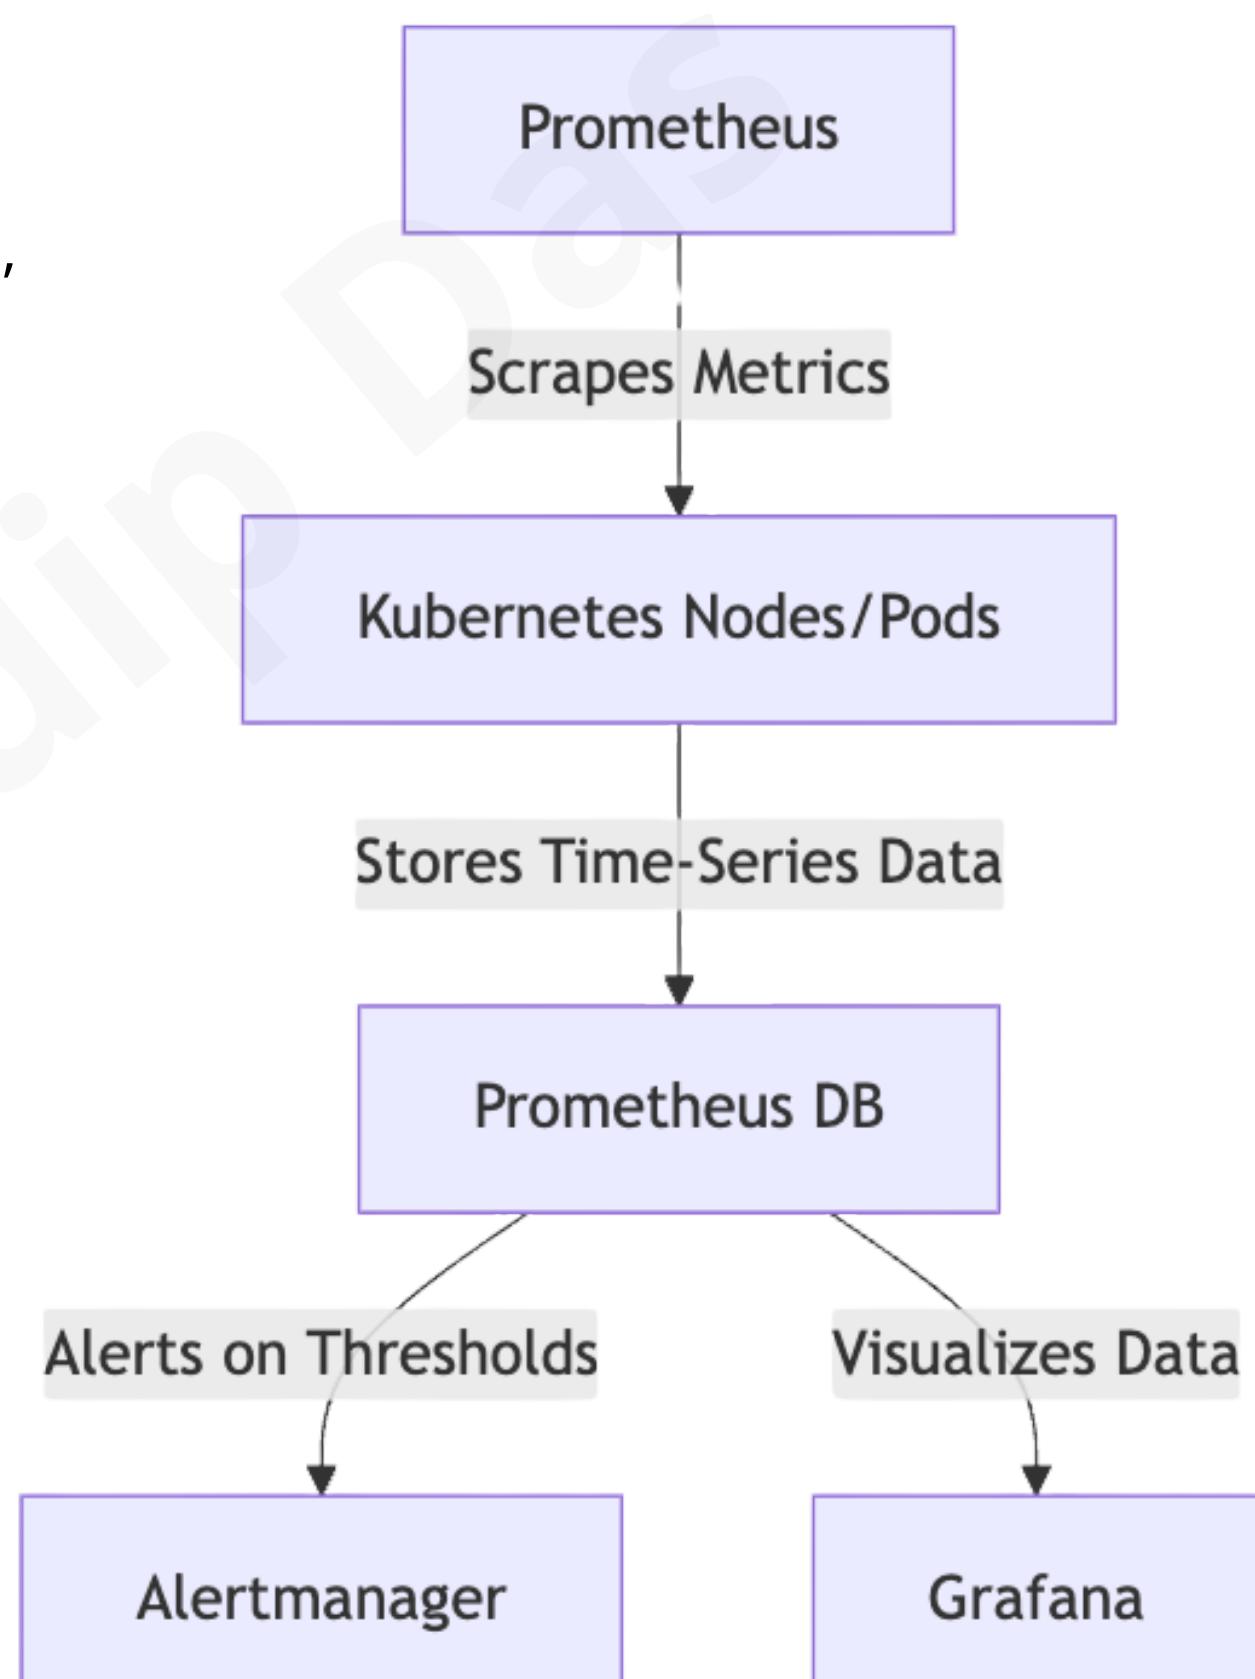
### Install Prometheus via Helm:

```
helm repo add prometheus-
community https://prometheus-
community.github.io/helm-charts
```

```
helm install prometheus
prometheus-community/kube-
prometheus-stack
```

### Define a Prometheus Alert Rule

```
apiVersion: monitoring.coreos.com/v1
kind: PrometheusRule
metadata:
  name: high-cpu-usage
spec:
  groups:
    - name: node-rules
      rules:
        - alert: HighCPUUsage
          expr: node_cpu_seconds_total > 80
          for: 5m
          labels:
            severity: critical
          annotations:
            summary: "High CPU usage detected"
```



### ✓ Benefits

- Real-time monitoring of CPU, memory, and requests.
- Alerts via Alertmanager (Slack, Email, PagerDuty).
- Grafana dashboards for visualization.



@Sandip Das

# Monitoring, Logging, and Observability

## EFK Stack: Centralized Logging

EFK (Elasticsearch, Fluentd, Kibana) enables log collection, storage, and visualization.

Install EFK via Helm:

```
helm repo add elastic https://helm.elastic.co
helm install elasticsearch elastic/elasticsearch
helm install kibana elastic/kibana
helm install fluentd stable/fluentd
```

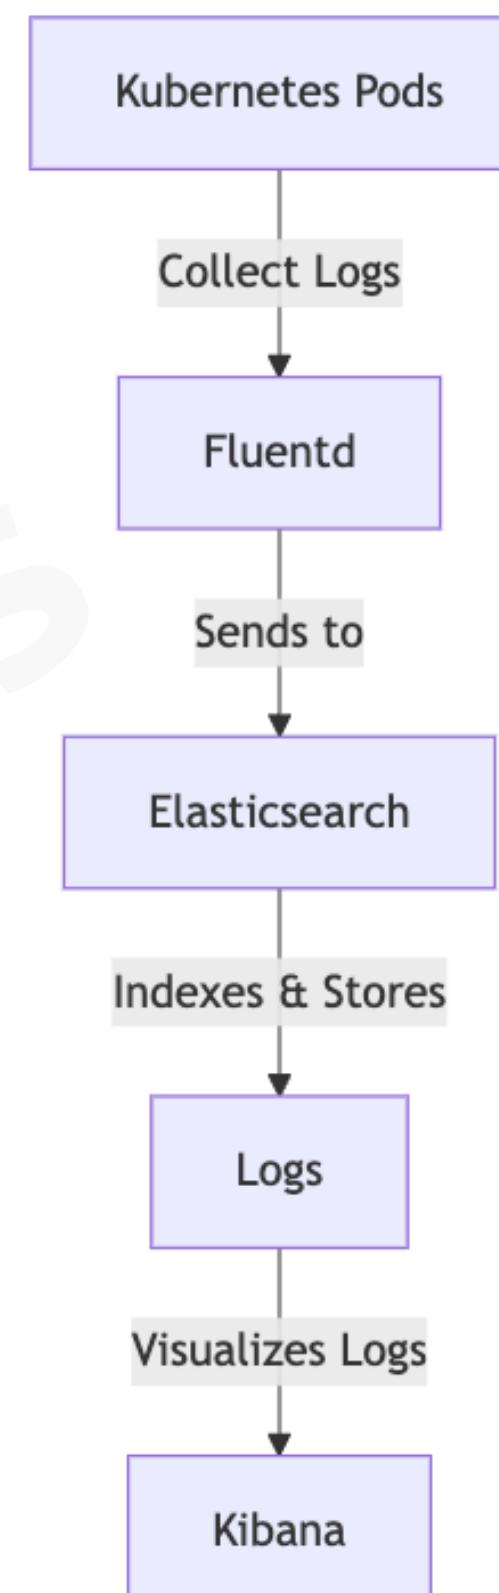
### Fluentd Configuration (Logging Pods)

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: fluentd-config
data:
  fluent.conf: |
    <source>
      @type tail
      path /var/log/containers/*.log
      pos_file /var/log/td-agent/containers.log.pos
      tag kubernetes.*
      format json
    </source>

    <match kubernetes.*>
      @type elasticsearch
      host elasticsearch
      port 9200
      logstash_format true
    </match>
```

#### Benefits

- Aggregates logs from all Kubernetes Pods.
- Elasticsearch stores logs, enabling fast searches.
- Kibana visualizes logs, helping with debugging.



# Monitoring, Logging, and Observability

## Distributed Tracing: Jaeger & OpenTelemetry

Tracing helps track requests as they move through microservices.

Install Jaeger via Helm:

```
helm repo add jaegertracing
https://jaegertracing.github.io/helm-charts
helm install jaeger jaegertracing/jaeger
```

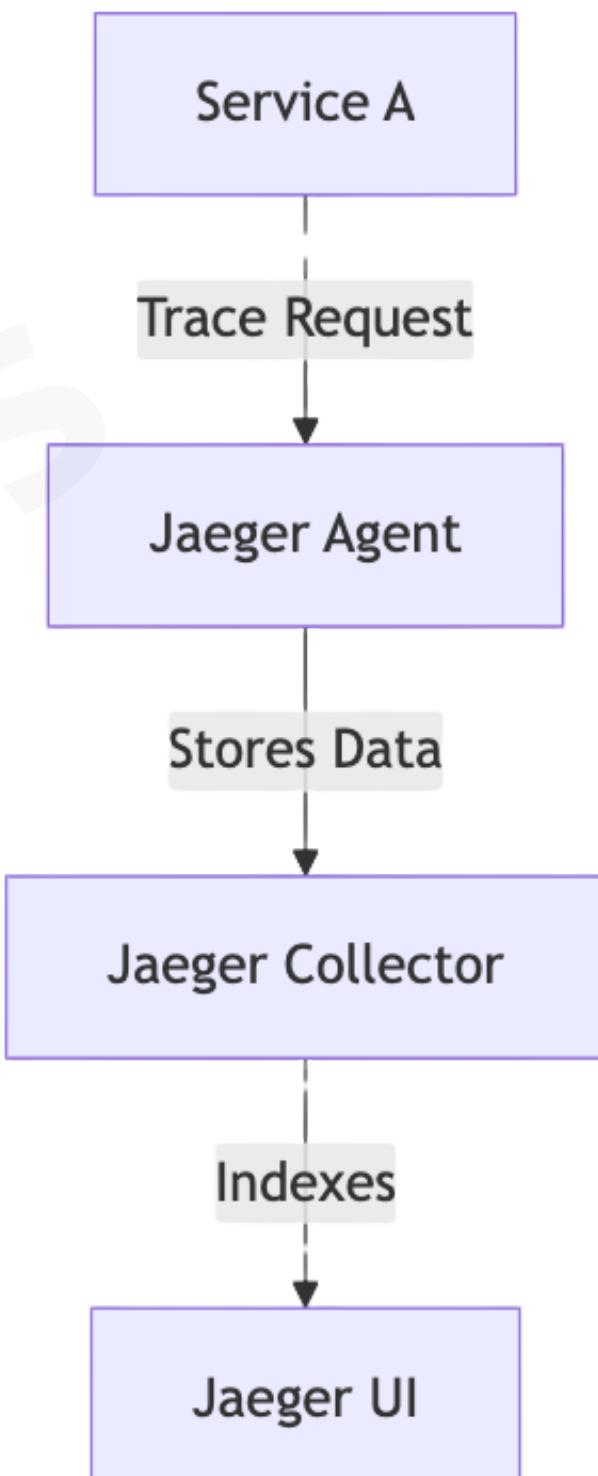
### Tracing Requests in Python (OpenTelemetry)

```
from opentelemetry import trace
from opentelemetry.exporter.jaeger.thrift import JaegerExporter
from opentelemetry.sdk.trace import TracerProvider
from opentelemetry.sdk.trace.export import SimpleSpanProcessor

# Setup Tracer
trace.set_tracer_provider(TracerProvider())
tracer = trace.get_tracer(__name__)

# Export to Jaeger
jaeger_exporter = JaegerExporter(agent_host_name="jaeger", agent_port=6831)
trace.get_tracer_provider().add_span_processor(SimpleSpanProcessor(jaeger_exporter))

# Create a span
with tracer.start_as_current_span("operation"):
    print("Tracing request in Jaeger")
```



### ✓ Benefits

- Identifies latency issues in microservices.
- Shows dependencies between services.
- Helps with debugging slow requests.



# CI/CD Pipelines with Kubernetes

CI/CD (Continuous Integration and Continuous Deployment) automates application building, testing, and deployment in Kubernetes. Now we will cover:

- ✓ GitOps with Argo CD → Declarative Kubernetes deployments
- ✓ Jenkins Pipelines → Automating build, test, and deployment
- ✓ Helm Charts → Packaging and managing Kubernetes releases

## GitOps with Argo CD: Declarative Deployments

GitOps uses Git as the single source of truth for Kubernetes configurations, ensuring declarative deployments.

### Install Argo CD:

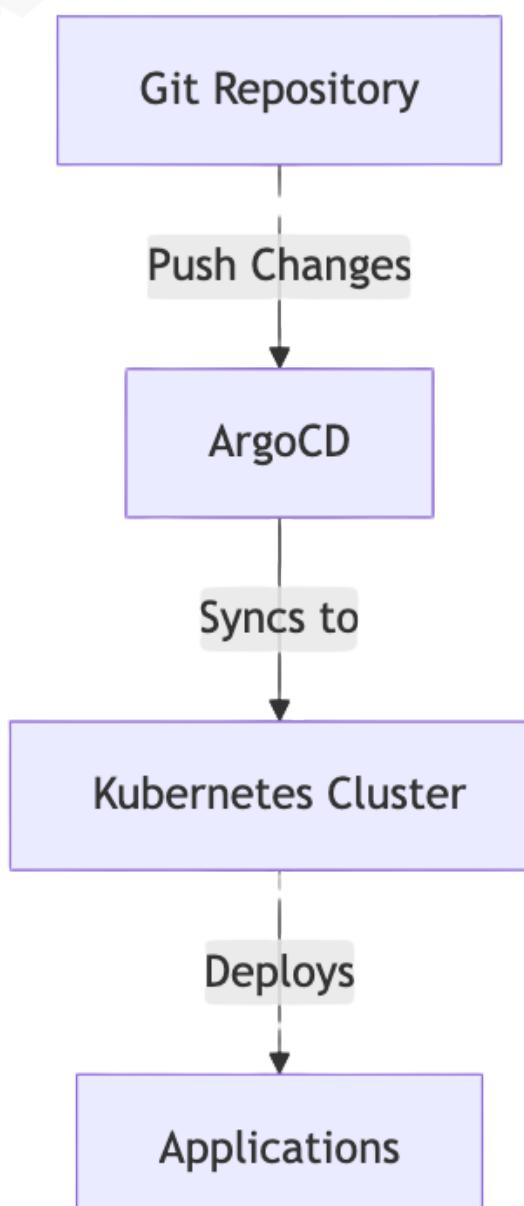
```
kubectl create namespace argocd  
kubectl apply -n argocd -f  
https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml
```

### Expose Argo CD Dashboard

```
kubectl port-forward svc/argocd-server -n argocd 8080:443
```

### Deploy an Application with Argo CD

```
apiVersion: argoproj.io/v1alpha1  
kind: Application  
metadata:  
  name: my-app  
  namespace: argocd  
spec:  
  destination:  
    namespace: my-app  
    server: https://kubernetes.default.svc  
  source:  
    repoURL: https://github.com/my-org/my-repo.git  
    targetRevision: main  
    path: my-app  
  syncPolicy:  
    automated:  
      prune: true  
      selfHeal: true
```



# CI/CD Pipelines with Kubernetes

## Jenkins Pipelines

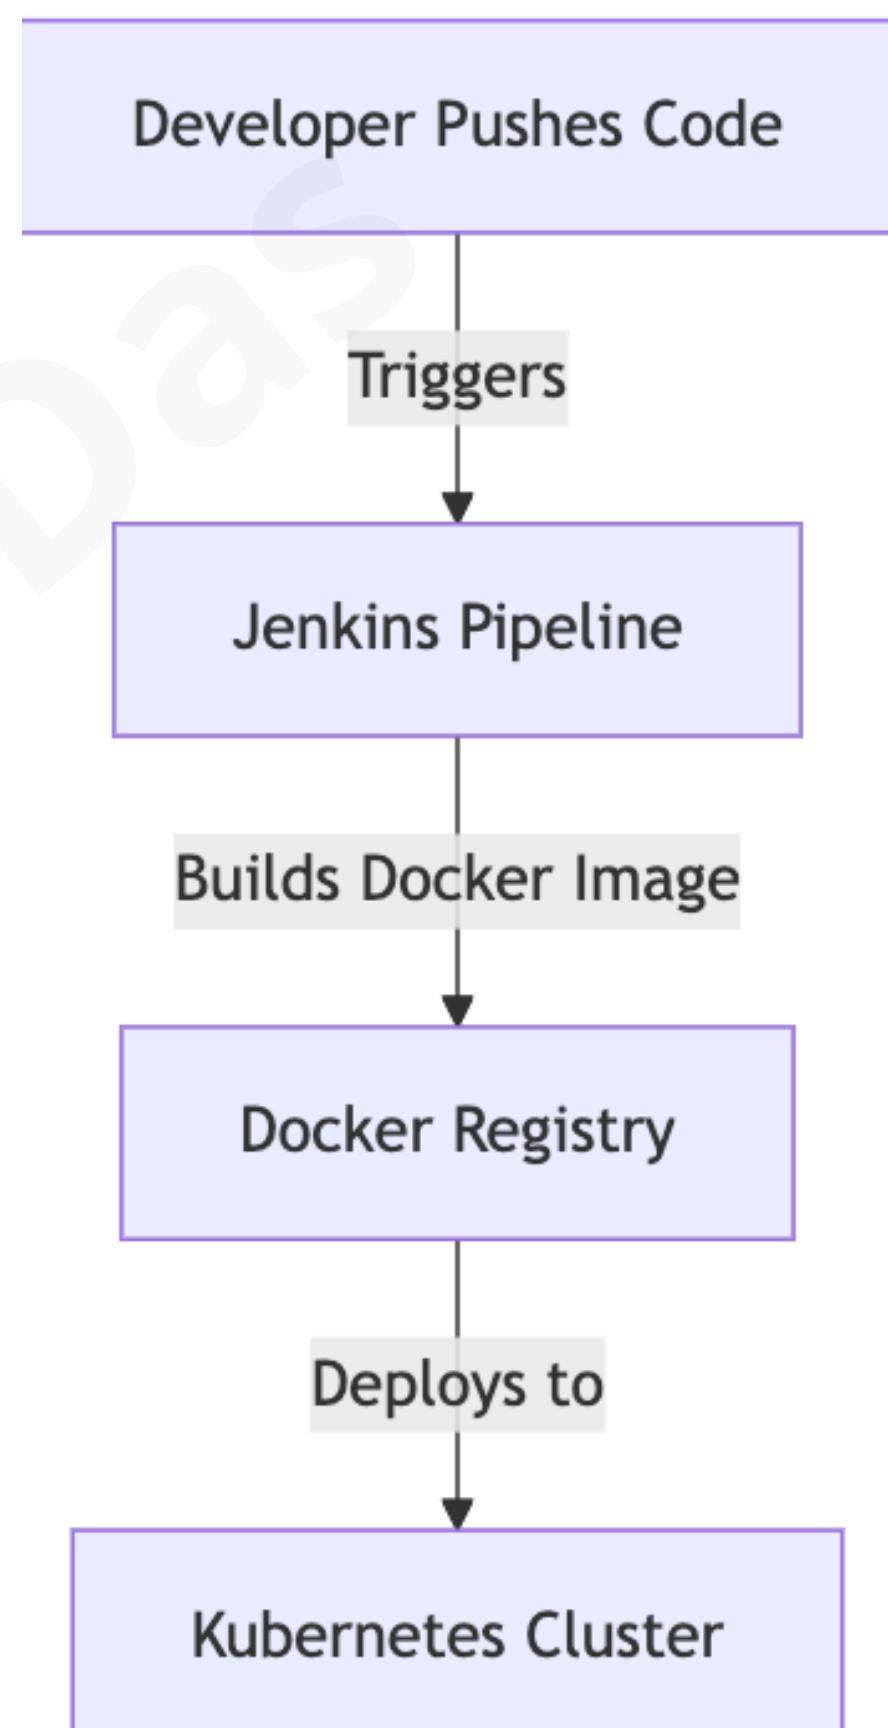
Jenkins automates building, testing, and deploying containerized applications.

### Install Jenkins on Kubernetes:

```
helm repo add jenkins https://charts.jenkins.io
helm install jenkins jenkins/jenkins
```

### Example: Jenkinsfile for Kubernetes CI/CD

```
pipeline {
    agent any
    environment {
        IMAGE_NAME = "my-app"
        IMAGE_TAG = "v1.0"
        REGISTRY = "docker.io/myrepo"
    }
    stages {
        stage('Build') {
            steps {
                sh 'docker build -t
$REGISTRY/$IMAGE_NAME:$IMAGE_TAG .'
            }
        }
        stage('Push to Registry') {
            steps {
                withDockerRegistry([credentialsId: 'docker-
credentials', url: 'https://index.docker.io/v1/']) {
                    sh 'docker push
$REGISTRY/$IMAGE_NAME:$IMAGE_TAG'
                }
            }
        }
        stage('Deploy to Kubernetes') {
            steps {
                sh 'kubectl apply -f k8s/deployment.yaml'
            }
        }
    }
}
```



# CI/CD Pipelines with Kubernetes

## Helm Charts: Managing Kubernetes Releases

Helm packages Kubernetes applications into reusable charts for easy deployment.

### Create a Helm Chart

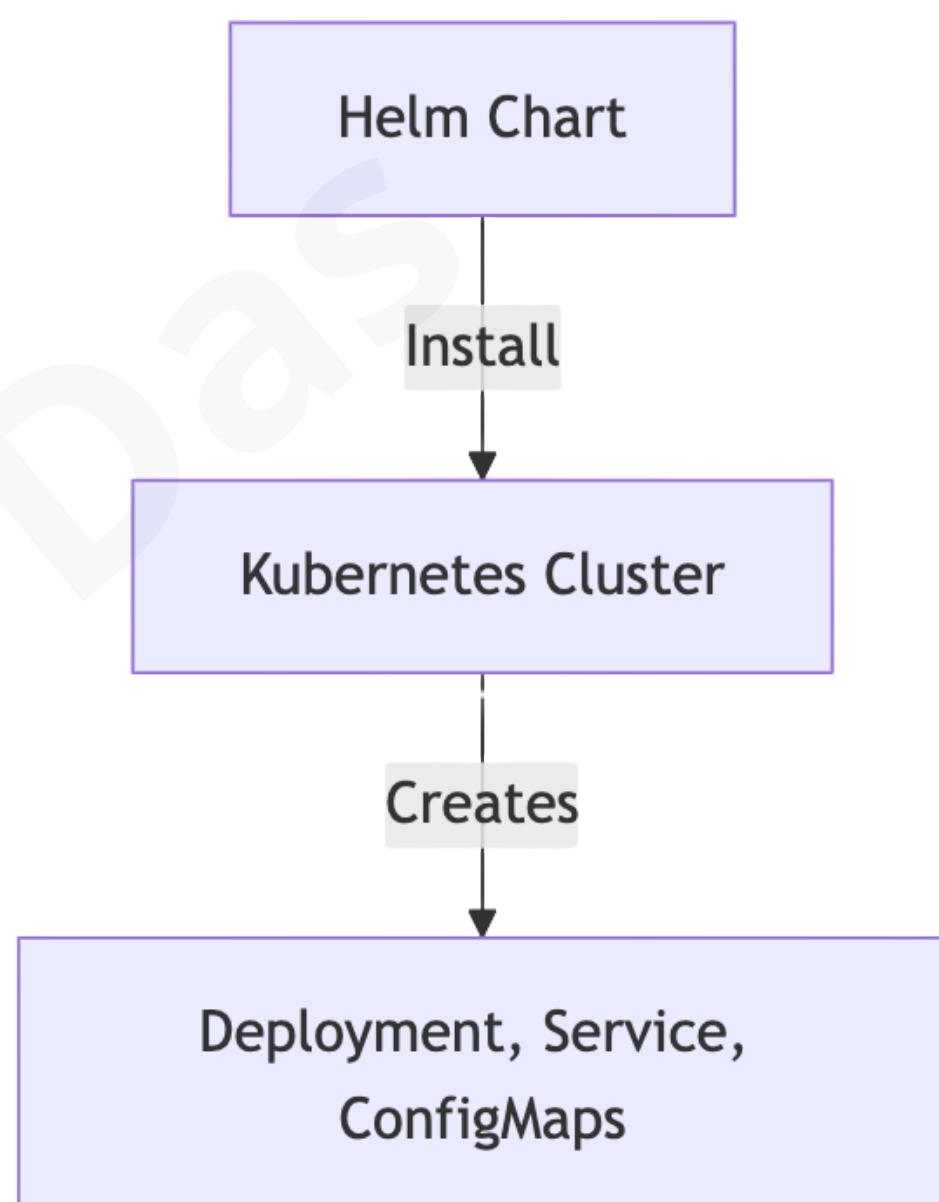
```
helm create my-app
```

### Deploy an Application with Helm

```
helm install my-app ./my-app
```

### Example: Helm values.yaml

```
replicaCount: 2
image:
  repository: nginx
  tag: latest
service:
  type: LoadBalancer
  port: 80
```



# What Next You May Explore?

## Serverless with Knative: Event-Driven Scaling

Knative brings serverless capabilities to Kubernetes, allowing applications to scale down to zero when idle.

### Key Features:

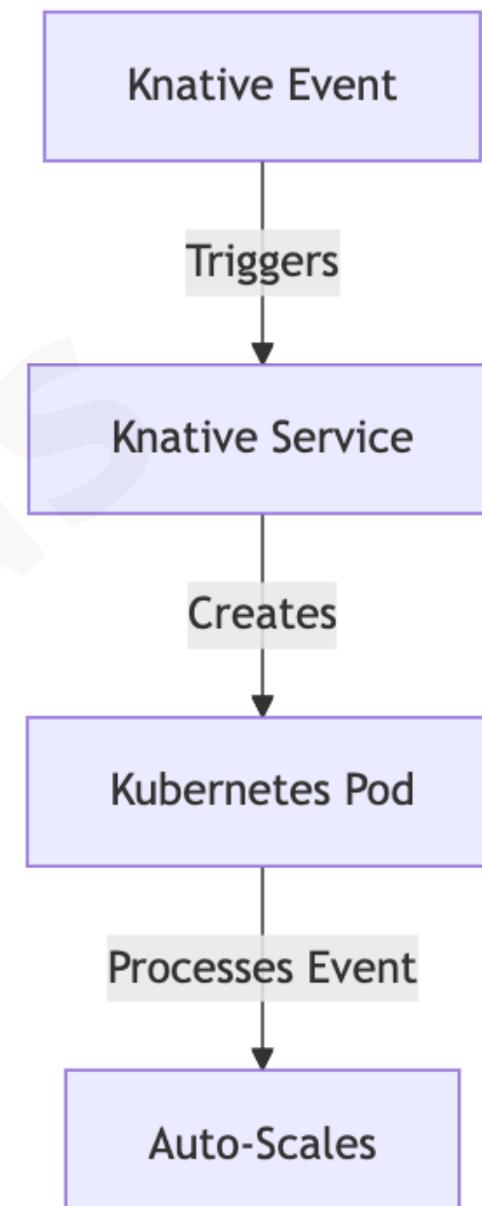
- Event-driven scaling → Pods scale up/down based on requests
- Works with Kubernetes → Deploy serverless workloads without extra infrastructure
- Cold start optimization → Efficient resource usage

### Installing Knative

```
kubectl apply -f  
https://github.com/knative/serving/releases/latest/download/serving-crds.yaml  
kubectl apply -f  
https://github.com/knative/serving/releases/latest/download/serving-core.yaml
```

### Example: Deploy a Serverless App

```
apiVersion: serving.knative.dev/v1  
kind: Service  
metadata:  
  name: hello-knative  
spec:  
  template:  
    spec:  
      containers:  
        - image: gcr.io/knative-samples/helloworld-go
```



# What Next You May Explore?

## Service Mesh with Istio: Advanced Traffic Management

Istio enhances Kubernetes networking with traffic routing, security, and observability.

### Key Features:

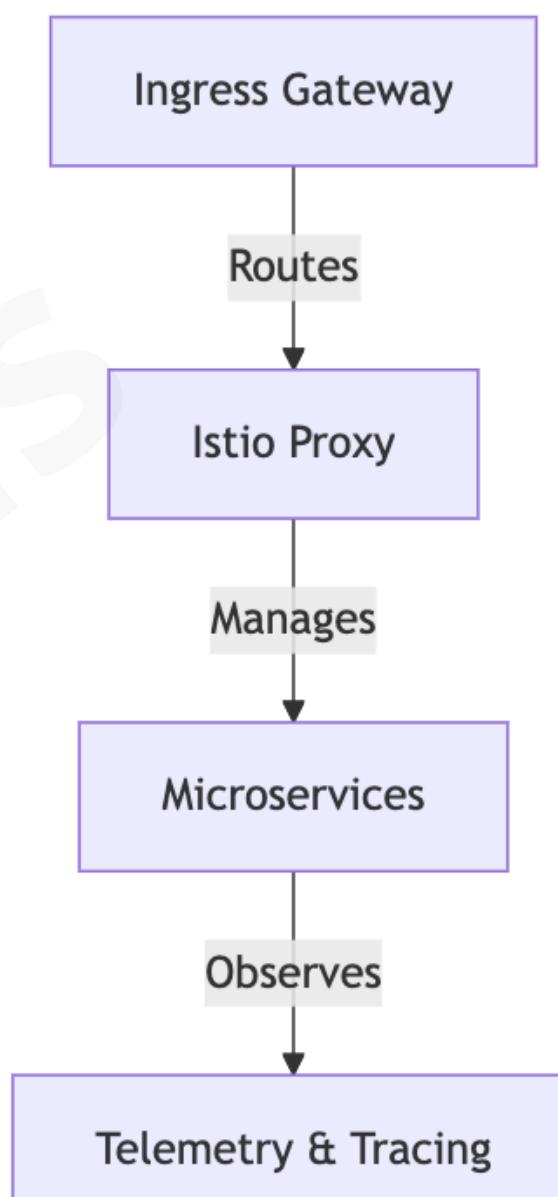
- ✓ Traffic management → A/B testing, canary releases
- ✓ Security → Mutual TLS (mTLS) encryption
- ✓ Observability → Built-in tracing and monitoring

### Install Istio

```
curl -L https://istio.io/downloadIstio | sh -
cd istio-*
export PATH=$PWD/bin:$PATH
istioctl install --set profile=demo -y
```

### Example: Canary Deployment

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: my-app
spec:
  hosts:
  - my-app.example.com
  http:
  - route:
    - destination:
        host: my-app
        subset: v1
        weight: 80
    - destination:
        host: my-app
        subset: v2
        weight: 20
```



# What Next You May Explore?

## Edge Computing with K3s & MicroK8s

Edge computing requires lightweight Kubernetes distributions optimized for low-resource environments.

### K3s (Lightweight Kubernetes)

- Designed for IoT & Edge
- Single binary, less than 100MB
- Faster startup than full Kubernetes

Install K3s on a Raspberry Pi

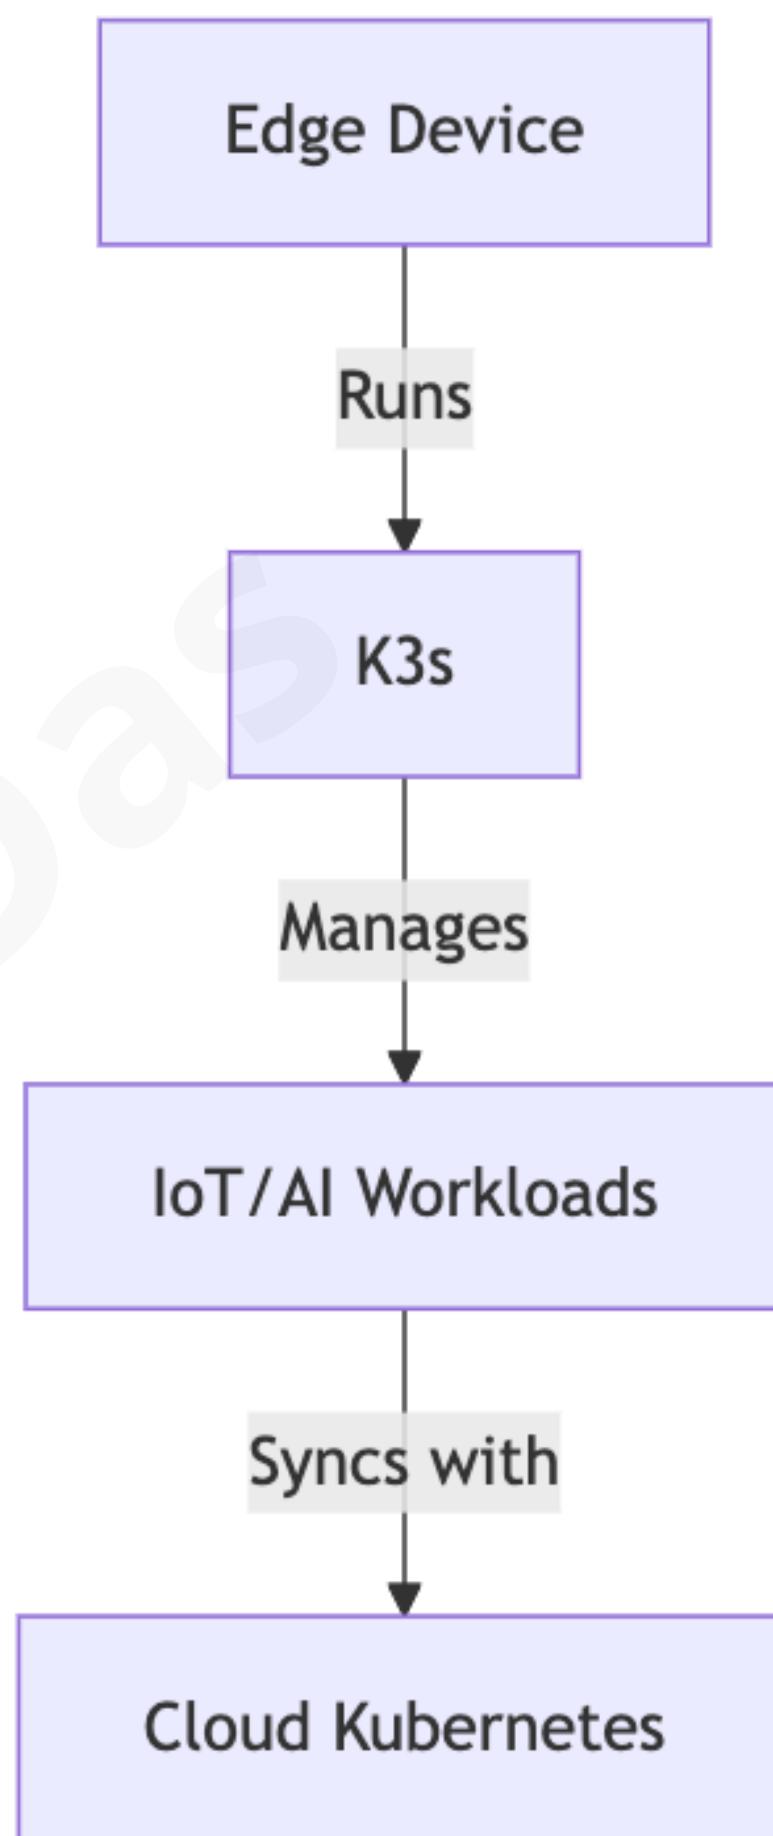
```
curl -sfL https://get.k3s.io | sh -
```

### MicroK8s (Canonical's Lightweight K8s)

- Optimized for local development & small clusters
- Snap-based install for easy upgrades
- Includes built-in storage & networking

Install MicroK8s

```
sudo snap install microk8s --classic
```



@Sandip Das

# GET SET GO

Keep Learning, Keep Building! 🚀

**Congratulations on completing Kubernetes Handbook - From Beginner to Expert!**

This is just the beginning of your journey in the world of Kubernetes and cloud-native technologies.

The tech landscape is ever-evolving, and the best way to stay ahead is to keep learning, experimenting, and sharing knowledge with others. Whether you are deploying your first containerized application or managing large-scale clusters, remember: every expert was once a beginner.

## 💡 What's next?

- Apply what you've learned by working on real-world projects.
- Contribute to open-source communities and collaborate with fellow engineers.
- Stay updated with the latest in Kubernetes and DevOps through blogs, conferences, and forums.
- Teach and mentor others – knowledge grows when shared!

If this book helped you, let me know! Connect with me on [LinkedIn](#) or follow my work at [LearnXOps](#).

Happy learning, and may your clusters always be healthy! 🚀🔥

- [Sandip Das](#)



@Sandip Das